

09 Amazon Fine Food Reviews Analysis_RF

April 28, 2019

1 Amazon Fine Food Reviews Analysis

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews>

EDA: <https://nycdatascience.com/blog/student-works/amazon-fine-foods-visualization/>

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454 Number of users: 256,059 Number of products: 74,258 Timespan: Oct 1999 - Oct 2012 Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

Objective: Given a review, determine whether the review is positive (rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative? [Ans] We could use Score/Rating. A rating of 4 or 5 can be considered as a positive review. A rating of 1 or 2 can be considered as negative one. A review of rating 3 is considered neutral and such reviews are ignored from our analysis. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

2 [1]. Reading Data

2.1 [1.1] Loading the data

The dataset is available in two forms 1. .csv file 2. SQLite Database

In order to load the data, We have used the SQLITE dataset as it is easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to "positive". Otherwise, it will be set to "negative".

```
In [1]: import warnings
        warnings.filterwarnings("ignore")

        from sklearn.metrics import roc_auc_score
        import sqlite3
        import pandas as pd
        import numpy as np
        import nltk
        import string
        import matplotlib.pyplot as plt
        import seaborn as sns
        !pip install -q PTable
        from prettytable import PrettyTable
        !pip3 install xgboost
        from sklearn.feature_extraction.text import TfidfTransformer
        from sklearn.feature_extraction.text import TfidfVectorizer
        from sklearn.model_selection import KFold
        from sklearn.metrics import roc_auc_score
        from sklearn.calibration import CalibratedClassifierCV
        from sklearn.model_selection import ParameterGrid
        from sklearn.preprocessing import StandardScaler
        from sklearn.feature_extraction.text import CountVectorizer
        from sklearn.metrics import confusion_matrix
        from sklearn.metrics import roc_curve, auc
        from nltk.stem.porter import PorterStemmer
        import re
        !pip3 install wordcloud
        # Tutorial about Python regular expressions: https://pymotw.com/2/re/
        import string
        from wordcloud import WordCloud
        from nltk.corpus import stopwords
        from nltk.stem import PorterStemmer
        from nltk.stem.wordnet import WordNetLemmatizer
        !pip install -q scikit-plot
        import scikitplot as skplt
        from gensim.models import Word2Vec
        from gensim.models import KeyedVectors
        import pickle
        #for finding nonzero elements in sparse matrix
        from scipy.sparse import find
        #for f1_Score
        from sklearn.metrics import f1_score
        #for displaying time
        from datetime import datetime
```

```

#for roc curve
import matplotlib.pyplot as plt
from itertools import cycle
from sklearn.ensemble import RandomForestClassifier
from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import label_binarize
from sklearn.multiclass import OneVsRestClassifier
from scipy.sparse import coo_matrix, hstack
from scipy import interp
from sklearn.metrics import classification_report
from sklearn.model_selection import GridSearchCV
from sklearn.tree import DecisionTreeClassifier
from tqdm import tqdm
import os
import xgboost as xgb

```

You are using pip version 19.0.3, however version 19.1 is available.

You should consider upgrading via the 'python -m pip install --upgrade pip' command.

Requirement already satisfied: xgboost in c:\users\shubh\anaconda3\lib\site-packages (0.82)

Requirement already satisfied: scipy in c:\users\shubh\anaconda3\lib\site-packages (from xgboost)

Requirement already satisfied: numpy in c:\users\shubh\anaconda3\lib\site-packages (from xgboost)

You are using pip version 19.0.3, however version 19.1 is available.

You should consider upgrading via the 'python -m pip install --upgrade pip' command.

Requirement already satisfied: wordcloud in c:\users\shubh\anaconda3\lib\site-packages (1.5.0)

Requirement already satisfied: numpy>=1.6.1 in c:\users\shubh\anaconda3\lib\site-packages (from wordcloud)

Requirement already satisfied: pillow in c:\users\shubh\anaconda3\lib\site-packages (from wordcloud)

You are using pip version 19.0.3, however version 19.1 is available.

You should consider upgrading via the 'python -m pip install --upgrade pip' command.

You are using pip version 19.0.3, however version 19.1 is available.

You should consider upgrading via the 'python -m pip install --upgrade pip' command.

C:\Users\shubh\Anaconda3\lib\site-packages\gensim\utils.py:1212: UserWarning: detected Windows warnings.warn("detected Windows; aliasing chunkize to chunkize_serial")

In [2]: # using SQLite Table to read data.

```

#os.chdir("/content/drive/Colab Notebooks") #changing directory
con = sqlite3.connect('database.sqlite')

```

```

# filtering only positive and negative reviews i.e.
# not taking into consideration those reviews with Score=3
# SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000, will give top 500000 data points
# you can change the number to any other number based on your computing power

#filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000""", con)
# for tsne assignment you can take 5k data points

filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3""", con)

# Give reviews with Score>3 a positive rating(1), and reviews with a score<3 a negative rating(-1)
def partition(x):
    if x < 3:
        return 0
    else:
        return 1

#changing reviews with score less than 3 to be positive and vice-versa
actualScore = filtered_data['Score']
positiveNegative = actualScore.map(partition)
filtered_data['Score'] = positiveNegative
print("Number of data points in our data", filtered_data.shape)
filtered_data.head(3)

```

Number of data points in our data (525814, 10)

```

Out[2]:
   Id  ProductId  UserId  ProfileName \
0   1  B001E4KFG0  A3SGXH7AUHU8GW  delmartian
1   2  B00813GRG4  A1D87F6ZCVE5NK  dll pa
2   3  B000LQOCHO  ABXLMWJIXXAIN  Natalia Corres "Natalia Corres"

   HelpfulnessNumerator  HelpfulnessDenominator  Score  Time \
0                      1                      1      1  1303862400
1                      0                      0      0  1346976000
2                      1                      1      1  1219017600

   Summary  Text
0  Good Quality Dog Food  I have bought several of the Vitality canned d...
1    Not as Advertised  Product arrived labeled as Jumbo Salted Peanut...
2  "Delight" says it all  This is a confection that has been around a fe...

```

```

In [3]: display = pd.read_sql_query("""
SELECT UserId, ProductId, ProfileName, Time, Score, Text, COUNT(*)
FROM Reviews
GROUP BY UserId
HAVING COUNT(*)>1
""", con)

```

```

In [4]: print(display.shape)
display.head()

```

(80668, 7)

```
Out [4]:
```

	UserId	ProductId	ProfileName	Time	Score	\
0	#oc-R115TNMSPFT9I7	B007Y59HVM	Breyton	1331510400	2	
1	#oc-R11D9D7SHXIJB9	B005HG9ETO	Louis E. Emory "hoppy"	1342396800	5	
2	#oc-R11DNU2NBKQ23Z	B007Y59HVM	Kim Cieszykowski	1348531200	1	
3	#oc-R1105J5ZVQE25C	B005HG9ETO	Penguin Chick	1346889600	5	
4	#oc-R12KPBODL2B5ZD	B0070SBE1U	Christopher P. Presta	1348617600	1	

	Text	COUNT(*)
0	Overall its just OK when considering the price...	2
1	My wife has recurring extreme muscle spasms, u...	3
2	This coffee is horrible and unfortunately not ...	2
3	This will be the bottle that you grab from the...	3
4	I didnt like this coffee. Instead of telling y...	2

```
In [5]: display[display['UserId']=='AZY10LLTJ71NX']
```

```
Out [5]:
```

	UserId	ProductId	ProfileName	Time	\
80638	AZY10LLTJ71NX	B006P7E5ZI	undertheshrine "undertheshrine"	1334707200	

	Score	Text	COUNT(*)
80638	5	I was recommended to try green tea extract to ...	5

```
In [6]: display['COUNT(*)'].sum()
```

```
Out [6]: 393063
```

3 [2] Exploratory Data Analysis

3.1 [2.1] Data Cleaning: Deduplication

It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

```
In [7]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND UserId="AR5J8UI46CURR"
ORDER BY ProductID
""", con)
display.head()
```

```
Out [7]:
```

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	\
0	78445	B000HDL1RQ	AR5J8UI46CURR	Geetha Krishnan	2	
1	138317	B000HDOPYC	AR5J8UI46CURR	Geetha Krishnan	2	
2	138277	B000HDOPYM	AR5J8UI46CURR	Geetha Krishnan	2	

3	73791	B000HDOPZG	AR5J8UI46CURR	Geetha Krishnan	2
4	155049	B000PAQ75C	AR5J8UI46CURR	Geetha Krishnan	2

	HelpfulnessDenominator	Score	Time \
0	2	5	1199577600
1	2	5	1199577600
2	2	5	1199577600
3	2	5	1199577600
4	2	5	1199577600

	Summary \
0	LOACKER QUADRATINI VANILLA WAFERS
1	LOACKER QUADRATINI VANILLA WAFERS
2	LOACKER QUADRATINI VANILLA WAFERS
3	LOACKER QUADRATINI VANILLA WAFERS
4	LOACKER QUADRATINI VANILLA WAFERS

	Text
0	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
1	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
2	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
3	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
4	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...

As it can be seen above that same user has multiple reviews with same values for HelpfulnessNumerator, HelpfulnessDenominator, Score, Time, Summary and Text and on doing analysis it was found that ProductId=B000HDOPZG was Loacker Quadratini Vanilla Wafer Cookies, 8.82-Ounce Packages (Pack of 8) ProductId=B000HDL1RQ was Loacker Quadratini Lemon Wafer Cookies, 8.82-Ounce Packages (Pack of 8) and so on

It was inferred after analysis that reviews with same parameters other than ProductId belonged to the same product just having different flavour or quantity. Hence in order to reduce redundancy it was decided to eliminate the rows having same parameters.

The method used for the same was that we first sort the data according to ProductId and then just keep the first similar product review and delete the others. for eg. in the above just the review for ProductId=B000HDL1RQ remains. This method ensures that there is only one representative for each product and deduplication without sorting would lead to possibility of different representatives still existing for the same product.

```
In [8]: #Sorting data according to ProductId in ascending order
sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True, inplace=False)
```

```
In [9]: #Deduplication of entries
final=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"}, keep='first')
final.shape
```

```
Out[9]: (364173, 10)
```

```
In [10]: #Checking to see how much % of data still remains
(final['Id'].size*1.0)/(filtered_data['Id'].size*1.0)*100
```

```
Out[10]: 69.25890143662969
```

Observation:- It was also seen that in two rows given below the value of HelpfulnessNumerator is greater than HelpfulnessDenominator which is not practically possible hence these two rows too are removed from calculations

```
In [11]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND Id=44737 OR Id=64422
ORDER BY ProductID
""", con)
```

```
display.head()
```

```
Out[11]:
```

	Id	ProductId	UserId	ProfileName	\
0	64422	B000MIDROQ	A161DK06JJMCYF	J. E. Stephens	"Jeanne"
1	44737	B001EQ55RW	A2V0I904FH7ABY		Ram

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	\
0		3	1	5	1224892800
1		3	2	4	1212883200

	Summary	\
0	Bought This for My Son at College	
1	Pure cocoa taste with crunchy almonds inside	

	Text
0	My son loves spaghetti so I didn't hesitate or...
1	It was almost a 'love at first bite' - the per...

```
In [12]: final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]
```

```
In [13]: #Before starting the next phase of preprocessing lets see the number of entries left
print(final.shape)
```

```
#How many positive and negative reviews are present in our dataset?
```

```
final['Score'].value_counts()
```

```
(364171, 10)
```

```
Out[13]: 1    307061
0     57110
Name: Score, dtype: int64
```

4 [3] Preprocessing

4.1 [3.1]. Preprocessing Review Text

Now that we have finished deduplication our data requires some preprocessing before we go on further with analysis and making the prediction model.

Hence in the Preprocessing phase we do the following in the order below:-

1. Begin by removing the html tags
2. Remove any punctuations or limited set of special characters like , or . or # etc.
3. Check if the word is made up of english letters and is not alpha-numeric
4. Check to see if the length of the word is greater than 2 (as it was researched that there is no adjective in 2-letters)
5. Convert the word to lowercase
6. Remove Stopwords
7. Finally Snowball Stemming the word (it was observed to be better than Porter Stemming)

After which we collect the words used to describe positive and negative reviews

```
In [14]: # printing some random reviews
sent_0 = final['Text'].values[0]
print(sent_0)
print("="*50)

sent_1000 = final['Text'].values[1000]
print(sent_1000)
print("="*50)

sent_1500 = final['Text'].values[1500]
print(sent_1500)
print("="*50)

sent_4900 = final['Text'].values[4900]
print(sent_4900)
print("="*50)
```

```
this witty little book makes my son laugh at loud. i recite it in the car as we're driving along
=====
I was really looking forward to these pods based on the reviews. Starbucks is good, but I prefer
=====
Great ingredients although, chicken should have been 1st rather than chicken broth, the only thing
=====
Can't do sugar. Have tried scores of SF Syrups. NONE of them can touch the excellence of this
=====
```

```
In [15]: # remove urls from text python: https://stackoverflow.com/a/40823105/4084039
sent_0 = re.sub(r"http\S+", "", sent_0)
sent_1000 = re.sub(r"http\S+", "", sent_1000)
```



```

sent_150 = re.sub(r"http\S+", "", sent_1500)
sent_4900 = re.sub(r"http\S+", "", sent_4900)

print(sent_0)

```

this witty little book makes my son laugh at loud. i recite it in the car as we're driving along

```

In [16]: # https://stackoverflow.com/questions/16206380/python-beautifulsoup-how-to-remove-all
from bs4 import BeautifulSoup

soup = BeautifulSoup(sent_0, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1000, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1500, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_4900, 'lxml')
text = soup.get_text()
print(text)

```

this witty little book makes my son laugh at loud. i recite it in the car as we're driving along

```

=====
I was really looking forward to these pods based on the reviews. Starbucks is good, but I prefer
=====
Great ingredients although, chicken should have been 1st rather than chicken broth, the only thing
=====
Can't do sugar. Have tried scores of SF Syrups. NONE of them can touch the excellence of this

```

```

In [17]: # https://stackoverflow.com/a/47091490/4084039
import re

def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"n't", " not", phrase)

```

```

phrase = re.sub(r"\ 're", " are", phrase)
phrase = re.sub(r"\ 's", " is", phrase)
phrase = re.sub(r"\ 'd", " would", phrase)
phrase = re.sub(r"\ 'll", " will", phrase)
phrase = re.sub(r"\ 't", " not", phrase)
phrase = re.sub(r"\ 've", " have", phrase)
phrase = re.sub(r"\ 'm", " am", phrase)
return phrase

```

```

In [18]: sent_1500 = decontracted(sent_1500)
print(sent_1500)
print("="*50)

```

Great ingredients although, chicken should have been 1st rather than chicken broth, the only thing
=====

```

In [19]: #remove words with numbers python: https://stackoverflow.com/a/18082370/4084039
sent_0 = re.sub("\S*\d\S*", "", sent_0).strip()
print(sent_0)

```

this witty little book makes my son laugh at loud. i recite it in the car as we're driving along

```

In [20]: #remove spacial character: https://stackoverflow.com/a/5843547/4084039
sent_1500 = re.sub('[^A-Za-z0-9]+', ' ', sent_1500)
print(sent_1500)

```

Great ingredients although chicken should have been 1st rather than chicken broth the only thing

```

In [21]: # https://gist.github.com/sebleier/554280
# we are removing the words from the stop words list: 'no', 'nor', 'not'
# <br /><br /> ==> after the above steps, we are getting "br br"
# we are including them into stop words list
# instead of <br /> if we have <br/> these tags would have reumoved in the 1st step

stopwords= set(['br', 'the', 'i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves',
'you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', 'him',
'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself',
'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', 'that',
'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'had',
'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'as',
'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'through',
'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off', 'over',
'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'any',
'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than', 'too',
's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've", 'n',
've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't",

```

```
"hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'mi',
"mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't",
'won', "won't", 'wouldn', "wouldn't"])
```

```
In [22]: # Combining all the above students
from tqdm import tqdm
preprocessed_reviews = []
# tqdm is for printing the status bar
for sentence in tqdm(final['Text'].values):
    sentence = re.sub(r"http\S+", "", sentence)
    sentence = BeautifulSoup(sentence, 'lxml').get_text()
    sentence = decontracted(sentence)
    sentence = re.sub("\S*\d\S*", "", sentence).strip()
    sentence = re.sub('[^A-Za-z]+', ' ', sentence)
    # https://gist.github.com/sebleier/554280
    sentence = ' '.join(e.lower() for e in sentence.split() if e.lower() not in stopwords)
    preprocessed_reviews.append(sentence.strip())
```

```
100%|| 364171/364171 [05:40<00:00, 1069.10it/s]
```

```
In [23]: preprocessed_reviews[1500]
```

```
Out[23]: 'great ingredients although chicken rather chicken broth thing not think belongs canola'
```

[3.2] Splitting the data

```
In [24]: final['Text'] = preprocessed_reviews
finalp = final[final.Score==1].sample(30000,random_state=2)
finaln = final[final.Score==0].sample(30000,random_state=2)
finalx = pd.concat([finalp,finaln],ignore_index=True)
finalx = finalx.sort_values('Time')
y = finalx.Score.values
X = finalx.Text.values
Xtr,Xtest,ytr,ytest = train_test_split(X,y,test_size = 0.3)
print(finalx.Score.value_counts())
print(Xtr.shape,Xtest.shape,ytr.shape,ytest.shape)
```

```
1    30000
0    30000
Name: Score, dtype: int64
(42000,) (18000,) (42000,) (18000,)
```

5 [4] Featurization

5.1 [4.1] BAG OF WORDS

```
In [25]: #BoW
count_bow = CountVectorizer(ngram_range=(1,2),min_df=10) #in scikit-learn
```

```

bow_vec_tr = count_bow.fit_transform(Xtr)
print("some feature names ", count_bow.get_feature_names()[:10])
print('='*50)
bow_vec_test = count_bow.transform(Xtest)

some feature names  ['aa', 'ability', 'able', 'able buy', 'able chew', 'able drink', 'able eat'
=====

```

5.2 [4.3] TF-IDF

```

In [26]: #tfidf
tf_idf_vect = TfidfVectorizer(ngram_range=(1,2), min_df=10)
tfidf_tr = tf_idf_vect.fit_transform(Xtr)
print("some sample features(unique words in the corpus)",tf_idf_vect.get_feature_names)
print('='*50)

tfidf_test = tf_idf_vect.transform(Xtest)
print("the type of count vectorizer ",type(tfidf_tr))
print("the shape of out text TFIDF vectorizer ",tfidf_tr.get_shape())
print("the number of unique words including both unigrams and bigrams ", tfidf_tr.get

some sample features(unique words in the corpus) ['aa', 'ability', 'able', 'able buy', 'able ch
=====
the type of count vectorizer  <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer  (42000, 25044)
the number of unique words including both unigrams and bigrams  25044

```

5.3 [4.4] Word2Vec

```

In [27]: # Train your own Word2Vec model using your own text corpus
i=0
list_of_sentence=[]
for sentence in Xtr:
    list_of_sentence.append(sentence.split())

In [28]: # Using Google News Word2Vectors

# in this project we are using a pretrained model by google
# its 3.3G file, once you load this into your memory
# it occupies ~9Gb, so please do this step only if you have >12G of ram
# we will provide a pickle file wich contains a dict ,
# and it contains all our courpus words as keys and model[word] as values
# To use this code-snippet, download "GoogleNews-vectors-negative300.bin"
# from https://drive.google.com/file/d/0B7XkCwpI5KDYNlNUTTlSS21pQmM/edit
# it's 1.9GB in size.

```

```
# http://kavita-ganesan.com/gensim-word2vec-tutorial-starter-code/#.W17SRFazZPY
# you can comment this whole cell
# or change these variable according to your need
```

```
is_your_ram_gt_16g=False
want_to_use_google_w2v = False
want_to_train_w2v = True
```

```
if want_to_train_w2v:
    # min_count = 5 considers only words that occurred atleast 5 times
    w2v_model=Word2Vec(list_of_sentence,min_count=5,size=50, workers=4)
    print(w2v_model.wv.most_similar('great'))
    print('='*50)
    print(w2v_model.wv.most_similar('worst'))
```

```
elif want_to_use_google_w2v and is_your_ram_gt_16g:
    if os.path.isfile('GoogleNews-vectors-negative300.bin'):
        w2v_model=KeyedVectors.load_word2vec_format('GoogleNews-vectors-negative300.b
        print(w2v_model.wv.most_similar('great'))
        print(w2v_model.wv.most_similar('worst'))
    else:
        print("you don't have gogole's word2vec file, keep want_to_train_w2v = True, t
```

```
[('awesome', 0.8741124868392944), ('fantastic', 0.8215399384498596), ('good', 0.80543911457061
=====
[(('nastiest', 0.8355489373207092), ('best', 0.7728778123855591), ('weakest', 0.708370089530944
```

```
In [29]: w2v_words = list(w2v_model.wv.vocab)
          print("number of words that occurred minimum 5 times ",len(w2v_words))
          print("sample words ", w2v_words[0:50])
```

number of words that occurred minimum 5 times 13054

sample words ['really', 'wanted', 'like', 'chocolate', 'cookies', 'reading', 'reviews', 'not'

5.4 [4.4.1] Converting text into vectors using Avg W2V, TFIDF-W2V

[4.4.1.1] Avg W2v

```
In [30]: # average Word2Vec for training data
```

```
i=0
list_of_sent_intr=[]
for sent in Xtr:
    list_of_sent_intr.append(sent.split())
```

```
# compute average word2vec for each review.
```

```
sent_vectors_intr = []; # the avg-w2v for each sentence/review is stored in this list
for sent in tqdm(list_of_sent_intr): # for each review/sentence
```

```

sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might need t
cnt_words = 0; # num of words with a valid vector in the sentence/review
for word in sent: # for each word in a review/sentence
    if word in w2v_words:
        vec = w2v_model.wv[word]
        sent_vec += vec
        cnt_words += 1
if cnt_words != 0:
    sent_vec /= cnt_words
sent_vectors_intr.append(sent_vec)
print(len(sent_vectors_intr))
print(len(sent_vectors_intr[0]))

# average Word2Vec for test data
i=0
list_of_sent_intest=[]
for sent in Xtest:
    list_of_sent_intest.append(sent.split())

# compute average word2vec for each review.
sent_vectors_intest = []; # the avg-w2v for each sentence/review is stored in this li
for sent in tqdm(list_of_sent_intest): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might need t
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    sent_vectors_intest.append(sent_vec)
print(len(sent_vectors_intest))
print(len(sent_vectors_intest[0]))

```

100%|| 42000/42000 [02:37<00:00, 266.68it/s]

42000
50

100%|| 18000/18000 [01:09<00:00, 260.16it/s]

18000
50

[4.4.1.2] TFIDF weighted W2v

```
In [31]: # S = ["abc def pqr", "def def def abc", "pqr pqr def"]
model = TfidfVectorizer(min_df=10)
tf_idf_matrix = model.fit_transform(Xtr)
model.transform(Xtest)
# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))

In [32]: # TF-IDF weighted Word2Vec
tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

tfidf_sent_vectors_intr = []; # the tfidf-w2v for each sentence/review is stored in t
row=0;
for sent in tqdm(list_of_sent_intr): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
#             tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
# to reduce the computation we are
# dictionary[word] = idf value of word in whole corpus
# sent.count(word) = tf value of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_sent_vectors_intr.append(sent_vec)
    row += 1

tfidf_sent_vectors_intest = []; # the tfidf-w2v for each sentence/review is stored in
row=0;
for sent in tqdm(list_of_sent_intest): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
#             tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
# to reduce the computation we are
# dictionary[word] = idf value of word in whole corpus
# sent.count(word) = tf value of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
```

```

        sent_vec /= weight_sum
    tfidf_sent_vectors_intest.append(sent_vec)
    row += 1

```

```

100%|| 42000/42000 [06:15<00:00, 111.78it/s]
100%|| 18000/18000 [01:48<00:00, 166.48it/s]

```

6 [5] Assignment 9: Random Forests

Apply Random Forests & GBDT on these feature sets

- SET 1:** Review text, preprocessed one converted into vectors

- SET 2:** Review text, preprocessed one converted into vectors

- SET 3:** Review text, preprocessed one converted into vectors

- SET 4:** Review text, preprocessed one converted into vectors

The hyper paramter tuning (Consider two hyperparameters: n_estimators & max_depth)

- Find the best hyper parameter which will give the maximum <https://www.appliedaicon.com>

- Find the best hyper paramter using k-fold cross validation or simple cross validation data

- Use gridsearch cv or randomsearch cv or you can also write your own for loops to do this task

Feature importance

- Get top 20 important features and represent them in a word cloud. Do this for BOW & TFIDF.

Feature engineering

- To increase the performance of your model, you can also experiment with with feature engineering

- Taking length of reviews as another feature.

- Considering some features from review summary as well.

Representation of results


```

<li>You need to plot the performance of model both on train data and cross validation data for
<img src='3d_plot.JPG' width=500px> with X-axis as <strong>n_estimators</strong>, Y-axis as <strong>(or)</strong></li>
<li>You need to plot the performance of model both on train data and cross validation data for
<img src='heat_map.JPG' width=300px> <a href='https://seaborn.pydata.org/generated/seaborn.heat_map.html'>https://seaborn.pydata.org/generated/seaborn.heat_map.html</a></li>
<li>You choose either of the plotting techniques out of 3d plot or heat map</li>
<li>Once after you found the best hyper parameter, you need to train your model with it, and find the best hyper parameter</li>
<img src='train_test_auc.JPG' width=300px></li>
<li>Along with plotting ROC curve, you need to print the <a href='https://www.appliedaicourse.com/roc-curve/'>https://www.appliedaicourse.com/roc-curve/</a></li>
<img src='confusion_matrix.png' width=300px></li>
</ul>
</li>
<br>
<li><strong>Conclusion</strong>
<ul>
<li>You need to summarize the results at the end of the notebook, summarize it in the table for
<img src='summary.JPG' width=400px>
</li>
</ul>

```

Note: Data Leakage

1. There will be an issue of data-leakage if you vectorize the entire data and then split it into train/cv/test.
2. To avoid the issue of data-leakage, make sure to split your data first and then vectorize it.
3. While vectorizing your data, apply the method fit_transform() on your train data, and apply the method transform() on cv/test data.
4. For more details please go through this link.

6.1 [5.1] Applying RF

Function for RF as well as XGBoost

```

In [33]: def rf(ft_train,ft_test,query):
    start = datetime.now()
    #Giving Parameters for tuning
    parameters = {'max_depth':[1, 5, 10, 50, 100, 500, 1000], 'n_estimators':[100, 500, 1000]}
    rf = RandomForestClassifier(n_jobs=-1)
    clf = GridSearchCV(rf, param_grid = parameters, scoring='roc_auc', cv=2,return_train_score=True)
    clf.fit(ft_train,ytr)

    results = clf.cv_results_
    train_score = results['mean_train_score']
    train_score_reshaped = train_score.reshape(7,4)
    test_score = results['mean_test_score']
    test_score_reshaped = test_score.reshape(7,4)
    max_depth=[1, 5, 10, 50, 100, 500, 1000]
    n_estimators=[100, 500, 700, 1000]

```

```

#Making into a Dataframe for Heatmaps
df_trainscore = pd.DataFrame(train_score_resaped,columns=n_estimators,index=max_depths)
df_testscore = pd.DataFrame(test_score_resaped,columns=n_estimators,index=max_depths)

#Getting Max Values
train_max_value = df_trainscore.values.max()
test_max_value = df_testscore.values.max()

#Finding location of the max values (row,column)
i1,j1= np.where(df_trainscore.values == train_max_value)
i2,j2 = np.where(df_testscore.values == test_max_value)
max_depth_train = list(df_trainscore.index[i1])[0]
n_est_train = list(df_trainscore.columns[j1])[0]
max_depth_test = list(df_testscore.index[i2])[0]
n_est_test = list(df_testscore.columns[j2])[0]

#Calculating Optimal Values
max_depth_optimal = int(np.median((max_depth_train,max_depth_test)))
n_est_optimal = int(np.median((n_est_train,n_est_test)))

#Plotting Heat Maps
fig, (ax1, ax2) =plt.subplots(1,2)
sns.heatmap(df_trainscore, annot = True, ax=ax1)
sns.heatmap(df_testscore, annot = True, ax=ax2)
ax1.set_title('Training plot')
ax1.set_xlabel('n_estimators')
ax1.set_ylabel('max_depth')
ax2.set_title('Validation plot')
ax2.set_xlabel('n_estimators')
ax2.set_ylabel('max_depth')
fig.show()

print('The maximum Train AUC is {} for {},{} . The max Validation AUC is {} for {}'.format(max_depth_train,n_est_train,max_depth_test,n_est_test))
print('Optimal parameters are max_depth = {} and n_estimators={} '.format(max_depth_optimal,n_est_optimal))
print("="*50)

#Training model with optimal parameters
model = RandomForestClassifier(max_depth=max_depth_optimal,n_estimators=n_est_optimal)
model.fit(ft_train,ytr)
pred_train = model.predict_proba(ft_train)
pred_test = model.predict_proba(ft_test)
p_train = model.predict(ft_train)
p_test = model.predict(ft_test)
f = model.feature_importances_

```

```
#Getting FPR AND TPR values for ROC Curve for train and test data
```

```
fpr = dict()
tpr = dict()
roc_auc = dict()
fpr,tpr,_ = roc_curve(ytr,pred_train[:,1])
roc_auc_train = roc_auc_score(ytr,pred_train[:,1])
fpr2 = dict()
tpr2 = dict()
roc_auc2 = dict()
fpr2,tpr2,_ = roc_curve(ytest,pred_test[:,1])
roc_auc_test = roc_auc_score(ytest,pred_test[:,1])
plt.figure()
plt.title(" ROC Curve")
plt.plot(fpr,tpr,'b',label='ROC curve for train data(area = %0.2f)' % roc_auc_train)
plt.plot(fpr2,tpr2,'r',label='ROC curve for test data(area = %0.2f)' % roc_auc_test)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend(loc="lower right")
plt.show()
#return max_depth_optimal,n_estimators_optimal
print('This is the ROC_AUC curve using optimal parameters with ROC_AUC of %0.2f for
print("="*50)

#For confusion matrix
print("Confusion Matrix for Train data")
skplt.metrics.plot_confusion_matrix(ytr,p_train)
print(classification_report(ytr,p_train))
print("="*50)
print("Confusion matrix for Test data")
skplt.metrics.plot_confusion_matrix(ytest,p_test)
print(classification_report(ytest,p_test))

print("Time taken to run this cell :", datetime.now() - start)
if query == 1:
    return f
```

In [34]: *#wordcloud*

```
def wcd(vector,w):
    features = vector.get_feature_names()
    top_features = pd.DataFrame(w,features)
    text = str(top_features[0].sort_values(ascending=False)[0:20])
    wordcloud = WordCloud(colormap="Oranges_r").generate(text)
    plt.imshow(wordcloud)
    plt.axis("off")
    plt.show()
```


6.1.1 [5.1.1] Applying Random Forests on BOW, SET 1


```
In [35]: w1 = rf(bow_vec_tr,bow_vec_test,1)
```

```
C:\Users\shubh\Anaconda3\lib\site-packages\matplotlib\figure.py:459: UserWarning: matplotlib is
"matplotlib is currently using a non-GUI backend, "
```

The maximum Train AUC is 0.9999970702909581 for 500,700 . The max Validation AUC is 0.9415635
Optimal parameters are max_depth = 300 and n_estimators=600

=====





This is the ROC_AUC curve using optimal parameters with ROC_AUC of 0.95 for test data



Confusion Matrix for Train data

	precision	recall	f1-score	support
0	1.00	1.00	1.00	21024
1	1.00	1.00	1.00	20976
avg / total	1.00	1.00	1.00	42000

Confusion matrix for Test data

	precision	recall	f1-score	support
0	0.87	0.89	0.88	8976
1	0.89	0.86	0.88	9024
avg / total	0.88	0.88	0.88	18000

Time taken to run this cell : 1:01:15.098462



6.1.2 [5.1.2] Wordcloud of top 20 important features from SET 1

```
In [36]: wcd(count_bow,w1)
```





6.1.3 [5.1.3] Applying Random Forests on TFIDF, SET 2

```
In [37]: w2 = rf(tfidf_tr,tfidf_test,1)
```

```
C:\Users\shubh\Anaconda3\lib\site-packages\matplotlib\figure.py:459: UserWarning: matplotlib is
"matplotlib is currently using a non-GUI backend, "
```

```
The maximum Train AUC is 0.9999970793612647 for 500,500 . The max Validation AUC is 0.9408070
Optimal parameters are max_depth = 500 and n_estimators=750
```

```
=====
```



This is the ROC_AUC curve using optimal parameters with ROC_AUC of 0.95 for test data

=====

Confusion Matrix for Train data


	precision	recall	f1-score	support
0	1.00	1.00	1.00	21024
1	1.00	1.00	1.00	20976
avg / total	1.00	1.00	1.00	42000


=====

Confusion matrix for Test data

	precision	recall	f1-score	support
0	0.87	0.89	0.88	8976
1	0.88	0.86	0.87	9024
avg / total	0.87	0.87	0.87	18000

Time taken to run this cell : 1:03:10.659817





6.1.4 [5.1.4] Wordcloud of top 20 important features from SET 2

In [38]: `wcd(tf_idf_vect,w2)`




6.1.5 [5.1.5] Applying Random Forests on AVG W2V, SET 3


```
In [39]: rf(sent_vectors_intr,sent_vectors_intest,0)
```

C:\Users\shubh\Anaconda3\lib\site-packages\matplotlib\figure.py:459: UserWarning: matplotlib is currently using a non-GUI backend, "

The maximum Train AUC is 0.999997090699148 for 1000,700 . The max Validation AUC is 0.9054132
Optimal parameters are max_depth = 750 and n_estimators=850

=====





This is the ROC_AUC curve using optimal parameters with ROC_AUC of 0.91 for test data



Confusion Matrix for Train data

	precision	recall	f1-score	support
0	1.00	1.00	1.00	21024
1	1.00	1.00	1.00	20976
avg / total	1.00	1.00	1.00	42000

Confusion matrix for Test data

	precision	recall	f1-score	support
0	0.83	0.84	0.83	8976
1	0.84	0.83	0.83	9024
avg / total	0.83	0.83	0.83	18000

Time taken to run this cell : 0:27:17.943669



6.1.6 [5.1.6] Applying Random Forests on TFIDF W2V, SET 4


```
In [40]: rf(tfidf_sent_vectors_intr,tfidf_sent_vectors_intest,0)
```

C:\Users\shubh\Anaconda3\lib\site-packages\matplotlib\figure.py:459: UserWarning: matplotlib is currently using a non-GUI backend, "

The maximum Train AUC is 0.9999970657558046 for 1000,700 . The max Validation AUC is 0.8849999999999999
Optimal parameters are max_depth = 1000 and n_estimators=850

=====





This is the ROC_AUC curve using optimal parameters with ROC_AUC of 0.89 for test data



Confusion Matrix for Train data

	precision	recall	f1-score	support
0	1.00	1.00	1.00	21024
1	1.00	1.00	1.00	20976
avg / total	1.00	1.00	1.00	42000

Confusion matrix for Test data

	precision	recall	f1-score	support
0	0.81	0.81	0.81	8976
1	0.81	0.81	0.81	9024
avg / total	0.81	0.81	0.81	18000

Time taken to run this cell : 0:27:11.411400



6.2 [5.2] Applying GBDT using XGBOOST

```
In [41]: def xg(ft_train,ft_test):
    start = datetime.now()
    #Giving Parameters for tuning
    parameters = {'max_depth':[1, 5, 10, 50, 100, 500, 1000], 'n_estimators':[100, 500, 1000]}
    xgboost = xgb.XGBClassifier(subsample=0.8, colsample_bytree=0.5, colsample_bylevel=0.5)
    clf = GridSearchCV(xgboost, param_grid = parameters, scoring='roc_auc', cv=2,return_train_score=True)
    clf.fit(ft_train,ytr)

    results = clf.cv_results_
    train_score = results['mean_train_score']
    train_score_reshaped = train_score.reshape(7,4)
    test_score = results['mean_test_score']
    test_score_reshaped = test_score.reshape(7,4)
    max_depth=[1, 5, 10, 50, 100, 500, 1000]
    n_estimators=[100, 500, 700, 1000]

    #Making into a Dataframe for Heatmaps
    df_trainscore = pd.DataFrame(train_score_reshaped,columns=n_estimators,index=max_depth)
    df_testscore = pd.DataFrame(test_score_reshaped,columns=n_estimators,index=max_depth)

    #Getting Max Values
    train_max_value = df_trainscore.values.max()
    test_max_value = df_testscore.values.max()

    #Finding location of the max values (row,column)
    i1,j1= np.where(df_trainscore.values == train_max_value)
    i2,j2 = np.where(df_testscore.values == test_max_value)
    max_depth_train = list(df_trainscore.index[i1])[0]
    n_est_train = list(df_trainscore.columns[j1])[0]
    max_depth_test = list(df_testscore.index[i2])[0]
    n_est_test = list(df_testscore.columns[j2])[0]

    #Calculating Optimal Values
    max_depth_optimal = int(np.median((max_depth_train,max_depth_test)))
    n_est_optimal = int(np.median((n_est_train,n_est_test)))

    #Plotting Heat Maps
    fig, (ax1, ax2) =plt.subplots(1,2)
    sns.heatmap(df_trainscore, annot = True, ax=ax1)
    sns.heatmap(df_testscore, annot = True, ax=ax2)
    ax1.set_title('Training plot')
    ax1.set_xlabel('n_estimators')
    ax1.set_ylabel('max_depth')
    ax2.set_title('Validation plot')
```

```

ax2.set_xlabel('n_estimators')
ax2.set_ylabel('max_depth')
fig.show()

print('The maximum Train AUC is {} for {},{} . The max Validation AUC is {} for {}'.format(max_train_auc, max_depth_optimal, n_estimators_optimal, max_validation_auc))
print('Optimal parameters are max_depth = {} and n_estimators={} '.format(max_depth_optimal, n_estimators_optimal))
print("="*50)

#Training model with optimal parameters
model = xgb.XGBClassifier(n_jobs=-1,max_depth=max_depth_optimal,n_estimators=n_estimators_optimal)
model.fit(ft_train,ytr)
pred_train = model.predict_proba(ft_train)
pred_test = model.predict_proba(ft_test)
p_train = model.predict(ft_train)
p_test = model.predict(ft_test)
f = model.feature_importances_

#Getting FPR AND TPR values for ROC Curve for train and test data

fpr = dict()
tpr = dict()
roc_auc = dict()
fpr,tpr,_ = roc_curve(ytr,pred_train[:,1])
roc_auc_train = roc_auc_score(ytr,pred_train[:,1])
fpr2 = dict()
tpr2 = dict()
roc_auc2 = dict()
fpr2,tpr2,_ = roc_curve(ytest,pred_test[:,1])
roc_auc_test = roc_auc_score(ytest,pred_test[:,1])
plt.figure()
plt.title(" ROC Curve")
plt.plot(fpr,tpr,'b',label='ROC curve for train data(area = %0.2f)' % roc_auc_train)
plt.plot(fpr2,tpr2,'r',label='ROC curve for test data(area = %0.2f)' % roc_auc_test)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend(loc="lower right")
plt.show()

#return max_depth_optimal,n_estimators_optimal
print('This is the ROC_AUC curve using optimal parameters with ROC_AUC of %0.2f for train data and %0.2f for test data' % (roc_auc_train, roc_auc_test))
print("="*50)

#For confusion matrix
print("Confusion Matrix for Train data")
skplt.metrics.plot_confusion_matrix(ytr,p_train)
print(classification_report(ytr,p_train))
print("="*50)
print("Confusion matrix for Test data")

```

```

skplt.metrics.plot_confusion_matrix(ytest,p_test)
print(classification_report(ytest,p_test))

print("Time taken to run this cell :", datetime.now() - start)

```

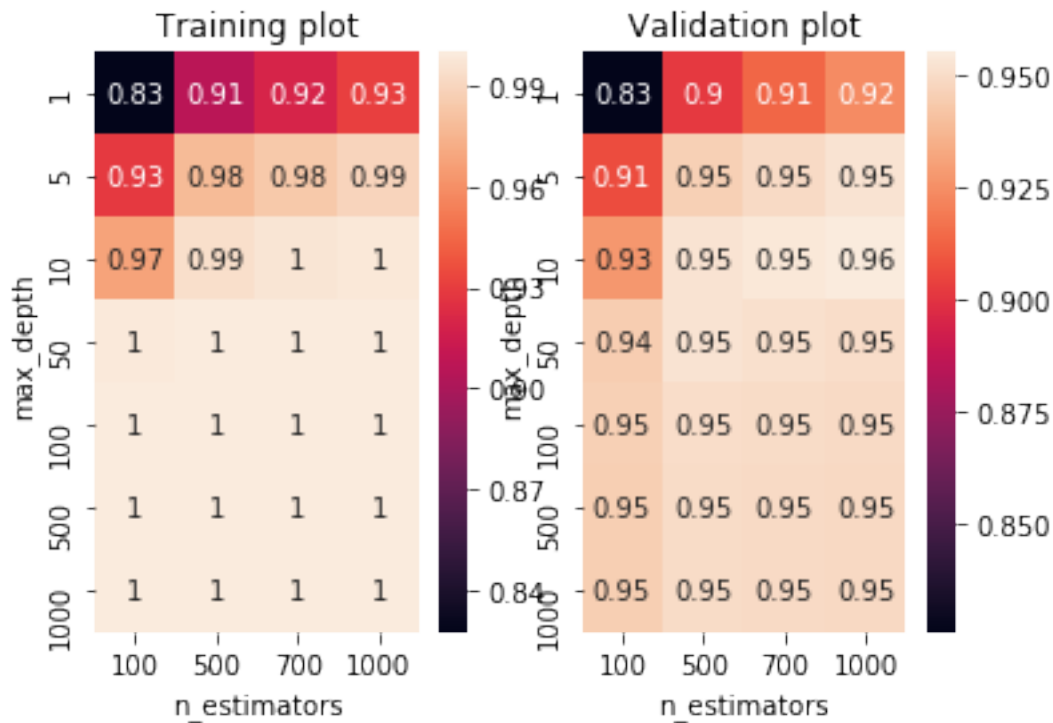
6.2.1 [5.2.1] Applying XGBOOST on BOW, SET 1


In [42]: `xg(bow_vec_tr,bow_vec_test)`

C:\Users\shubh\Anaconda3\lib\site-packages\matplotlib\figure.py:459: UserWarning: matplotlib is currently using a non-GUI backend, "

The maximum Train AUC is 0.9999968594063289 for 500,500 . The max Validation AUC is 0.9551075
Optimal parameters are max_depth = 255 and n_estimators=750

C:\Users\shubh\Anaconda3\lib\site-packages\sklearn\preprocessing\label.py:151: DeprecationWarning: if diff:
C:\Users\shubh\Anaconda3\lib\site-packages\sklearn\preprocessing\label.py:151: DeprecationWarning: if diff:





This is the ROC_AUC curve using optimal parameters with ROC_AUC of 0.96 for test data



Confusion Matrix for Train data

	precision	recall	f1-score	support
0	1.00	1.00	1.00	21024
1	1.00	1.00	1.00	20976
avg / total	1.00	1.00	1.00	42000

Confusion matrix for Test data

	precision	recall	f1-score	support
0	0.90	0.89	0.89	8976
1	0.89	0.90	0.89	9024
avg / total	0.89	0.89	0.89	18000

Time taken to run this cell : 3:22:02.227670



6.2.2 [5.2.2] Applying XGBOOST on TFIDF, SET 2

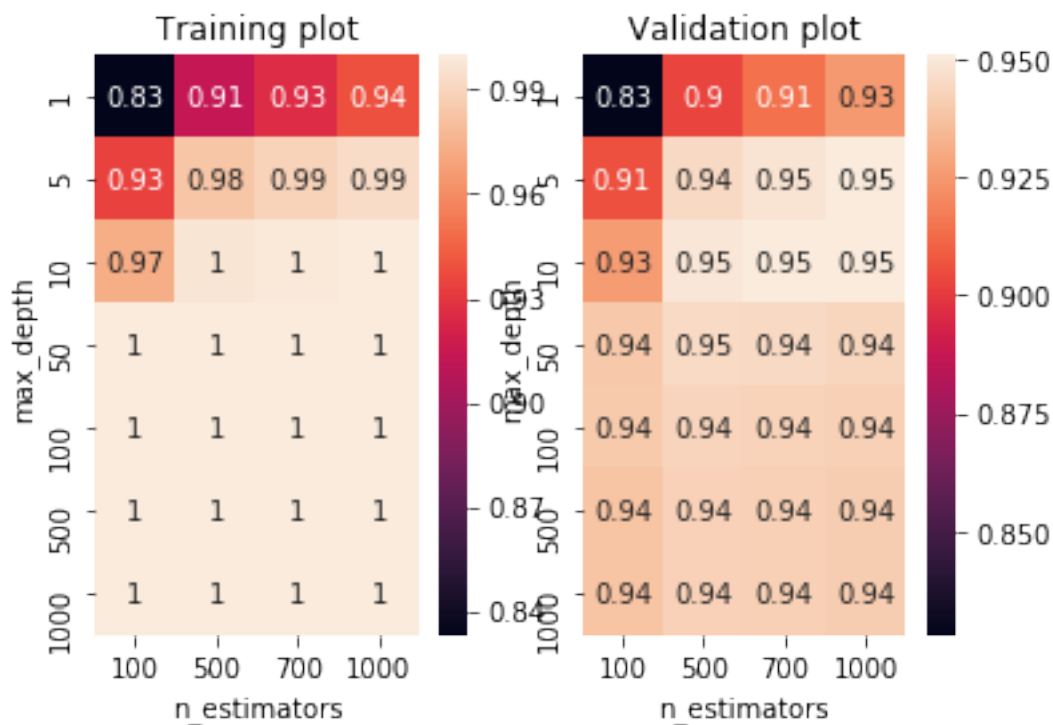
```
In [43]: xg(tfidf_tr,tfidf_test)
```


```
C:\Users\shubh\Anaconda3\lib\site-packages\matplotlib\figure.py:459: UserWarning: matplotlib is
"matplotlib is currently using a non-GUI backend, "
```

```
The maximum Train AUC is 0.9999969410390885 for 50,700 . The max Validation AUC is 0.95093143
Optimal parameters are max_depth = 27 and n_estimators=850
```

```
=====
```

```
C:\Users\shubh\Anaconda3\lib\site-packages\sklearn\preprocessing\label.py:151: DeprecationWarn
if diff:
C:\Users\shubh\Anaconda3\lib\site-packages\sklearn\preprocessing\label.py:151: DeprecationWarn
if diff:
```





This is the ROC_AUC curve using optimal parameters with ROC_AUC of 0.96 for test data



Confusion Matrix for Train data

	precision	recall	f1-score	support
0	1.00	1.00	1.00	21024
1	1.00	1.00	1.00	20976
avg / total	1.00	1.00	1.00	42000

Confusion matrix for Test data

	precision	recall	f1-score	support
0	0.89	0.89	0.89	8976
1	0.89	0.89	0.89	9024
avg / total	0.89	0.89	0.89	18000

Time taken to run this cell : 7:27:16.417076



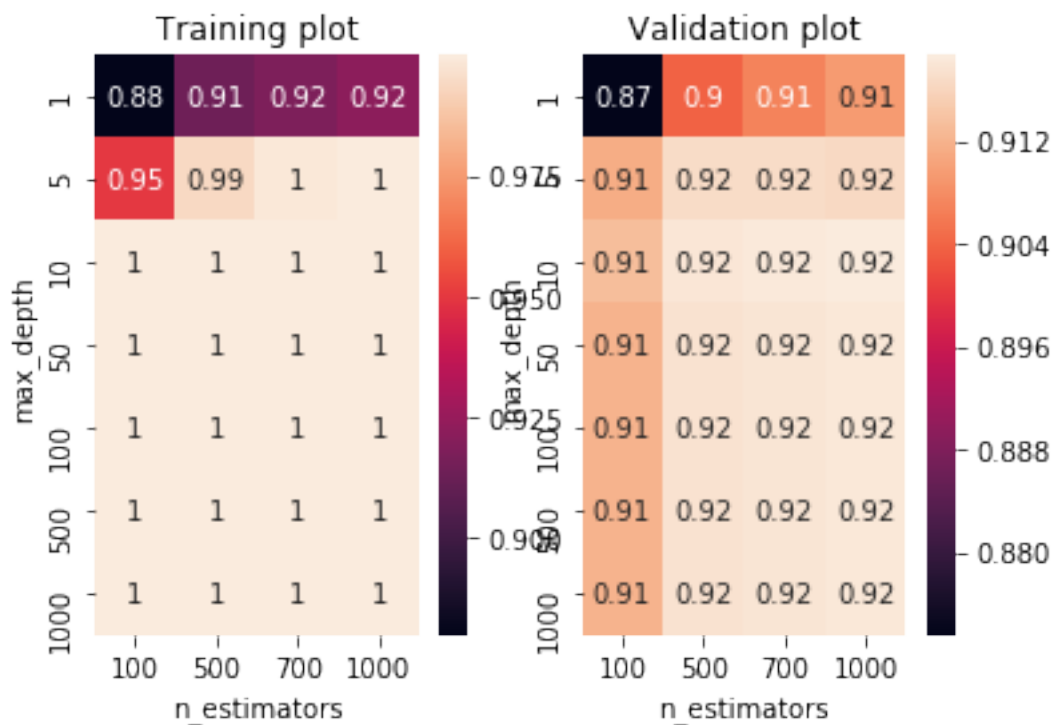
6.2.3 [5.2.3] Applying XGBOOST on AVG W2V, SET 3


```
In [44]: w2vtrain = np.array(sent_vectors_intr)
        w2vtest = np.array(sent_vectors_intest)
        xg(w2vtrain,w2vtest)
```

C:\Users\shubh\Anaconda3\lib\site-packages\matplotlib\figure.py:459: UserWarning: matplotlib is currently using a non-GUI backend, "

The maximum Train AUC is 0.9999971315155278 for 10,500 . The max Validation AUC is 0.91868332
Optimal parameters are max_depth = 10 and n_estimators=750

C:\Users\shubh\Anaconda3\lib\site-packages\sklearn\preprocessing\label.py:151: DeprecationWarning: if diff:
C:\Users\shubh\Anaconda3\lib\site-packages\sklearn\preprocessing\label.py:151: DeprecationWarning: if diff:





This is the ROC_AUC curve using optimal parameters with ROC_AUC of 0.92 for test data


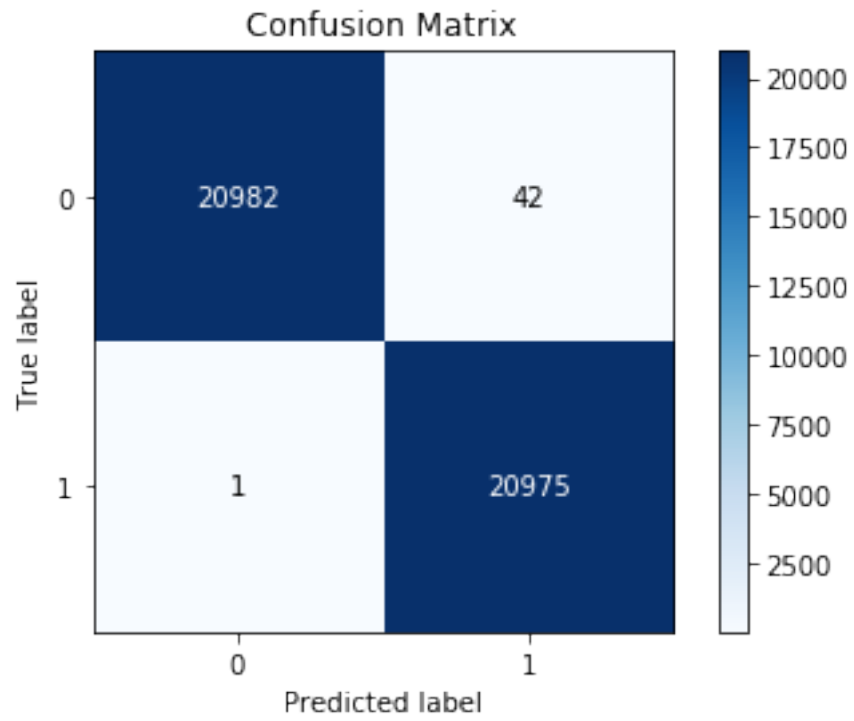
Confusion Matrix for Train data

	precision	recall	f1-score	support
0	1.00	1.00	1.00	21024
1	1.00	1.00	1.00	20976
avg / total	1.00	1.00	1.00	42000

Confusion matrix for Test data

	precision	recall	f1-score	support
0	0.84	0.85	0.84	8976
1	0.85	0.84	0.84	9024
avg / total	0.84	0.84	0.84	18000

Time taken to run this cell : 0:29:05.095060




6.2.4 [5.2.4] Applying XGBOOST on TFIDF W2V, SET 4


```
In [45]: tfw2vtrain = np.array(tfidf_sent_vectors_intr)
         tfw2vtest = np.array(tfidf_sent_vectors_intest)
         xg(tfw2vtrain,tfw2vtest)
```

C:\Users\shubh\Anaconda3\lib\site-packages\matplotlib\figure.py:459: UserWarning: matplotlib is currently using a non-GUI backend, "

The maximum Train AUC is 0.9999970793612647 for 10,500 . The max Validation AUC is 0.89935394
Optimal parameters are max_depth = 30 and n_estimators=750

C:\Users\shubh\Anaconda3\lib\site-packages\sklearn\preprocessing\label.py:151: DeprecationWarning: if diff:
C:\Users\shubh\Anaconda3\lib\site-packages\sklearn\preprocessing\label.py:151: DeprecationWarning: if diff:





This is the ROC_AUC curve using optimal parameters with ROC_AUC of 0.90 for test data



Confusion Matrix for Train data

	precision	recall	f1-score	support
0	1.00	1.00	1.00	21024
1	1.00	1.00	1.00	20976
avg / total	1.00	1.00	1.00	42000

Confusion matrix for Test data

	precision	recall	f1-score	support
0	0.82	0.82	0.82	8976
1	0.82	0.82	0.82	9024
avg / total	0.82	0.82	0.82	18000

Time taken to run this cell : 0:31:43.692104



7 [6] Conclusions

In [47]: x = PrettyTable()

```
x.field_names = ["Algorithm", "Vectorizer", "max_depth", "n_estimators", "AUC"]
x.add_row(["RandomForest", "BoW", 300, 600, 0.95])
x.add_row(["RandomForest", "Tfidf", 500, 750, 0.95])
x.add_row(["RandomForest", "Avg W2V", 750, 850, 0.91])
x.add_row(["RandomForest", "Tfidf weighted W2V", 1000, 850, 0.89])
x.add_row(["XGBoost", "BoW", 255, 750, 0.96])
x.add_row(["XGBoost", "Tfidf", 27, 850, 0.96])
x.add_row(["XGBoost", "Avg W2V", 10, 750, 0.92])
x.add_row(["XGBoost", "Tfidf weighted W2V", 30, 750, 0.90])
print(x)
```

Algorithm	Vectorizer	max_depth	n_estimators	AUC
RandomForest	BoW	300	600	0.95
RandomForest	Tfidf	500	750	0.95
RandomForest	Avg W2V	750	850	0.91
RandomForest	Tfidf weighted W2V	1000	850	0.89
XGBoost	BoW	255	750	0.96
XGBoost	Tfidf	27	850	0.96
XGBoost	Avg W2V	10	750	0.92
XGBoost	Tfidf weighted W2V	30	750	0.9

Conclusions: 1) Both Random Forest and XGBoost have a good accuracy but also have high time complexity. 2) These models can be trusted since they have high AUC. 3) Random Forest is based on bagging which tends to reduce variance without affecting Bias. GBDT is based on boosting which tends to reduce Bias without affecting Variance. However, GBDT based on XGBoost takes the best of both using GBDT and Random Forest using row sampling/col. sampling. Thus, GBDT using XGBoost has better AUC than RandomForest as it is evident from the above table.