

11 Amazon Fine Food Reviews Analysis_Truncated SVD

May 23, 2019

1 Amazon Fine Food Reviews Analysis

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews>

EDA: <https://nycdatascience.com/blog/student-works/amazon-fine-foods-visualization/>

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454 Number of users: 256,059 Number of products: 74,258 Timespan: Oct 1999 - Oct 2012 Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

Objective: Given a review, determine whether the review is positive (rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative? [Ans] We could use Score/Rating. A rating of 4 or 5 can be considered as a positive review. A rating of 1 or 2 can be considered as negative one. A review of rating 3 is considered neutral and such reviews are ignored from our analysis. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

2 [1]. Reading Data

2.1 [1.1] Loading the data

The dataset is available in two forms 1. .csv file 2. SQLite Database

In order to load the data, We have used the SQLITE dataset as it is easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to "positive". Otherwise, it will be set to "negative".

```
In [1]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os
from wordcloud import WordCloud, STOPWORDS
from sklearn.metrics.pairwise import cosine_similarity
```

```
C:\Users\shubh\Anaconda3\lib\site-packages\gensim\utils.py:1212: UserWarning: detected Windows
warnings.warn("detected Windows; aliasing chunkize to chunkize_serial")
```

```
In [2]: # using SQLite Table to read data.
con = sqlite3.connect('database.sqlite')
```

```

# filtering only positive and negative reviews i.e.
# not taking into consideration those reviews with Score=3
# SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000, will give top 500000 data points
# you can change the number to any other number based on your computing power

filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000
# for tsne assignment you can take 5k data points

#filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 5000

# Give reviews with Score>3 a positive rating(1), and reviews with a score<3 a negative rating(-1)
def partition(x):
    if x < 3:
        return 0
    else:
        return 1

#changing reviews with score less than 3 to be positive and vice-versa
actualScore = filtered_data['Score']
positiveNegative = actualScore.map(partition)
filtered_data['Score'] = positiveNegative
print("Number of data points in our data", filtered_data.shape)
filtered_data.head(3)

```

Number of data points in our data (500000, 10)

```

Out[2]:
   Id  ProductId  UserId  ProfileName \
0   1  B001E4KFG0  A3SGXH7AUHU8GW  delmartian
1   2  B00813GRG4  A1D87F6ZCVE5NK  dll pa
2   3  B000LQOCHO  ABXLMWJIXXAIN  Natalia Corres "Natalia Corres"

   HelpfulnessNumerator  HelpfulnessDenominator  Score  Time \
0                      1                      1      1  1303862400
1                      0                      0      0  1346976000
2                      1                      1      1  1219017600

   Summary  Text
0  Good Quality Dog Food  I have bought several of the Vitality canned d...
1  Not as Advertised  Product arrived labeled as Jumbo Salted Peanut...
2  "Delight" says it all  This is a confection that has been around a fe...

```

```

In [3]: display = pd.read_sql_query("""
SELECT UserId, ProductId, ProfileName, Time, Score, Text, COUNT(*)
FROM Reviews
GROUP BY UserId
HAVING COUNT(*)>1
""", con)

```

```

In [4]: print(display.shape)
display.head()

```

(80668, 7)

```
Out [4]:
```

	UserId	ProductId	ProfileName	Time	Score	\
0	#oc-R115TNMSPFT9I7	B007Y59HVM	Breyton	1331510400	2	
1	#oc-R11D9D7SHXIJB9	B005HG9ETO	Louis E. Emory "hoppy"	1342396800	5	
2	#oc-R11DNU2NBKQ23Z	B007Y59HVM	Kim Cieszykowski	1348531200	1	
3	#oc-R1105J5ZVQE25C	B005HG9ETO	Penguin Chick	1346889600	5	
4	#oc-R12KPBODL2B5ZD	B0070SBE1U	Christopher P. Presta	1348617600	1	

	Text	COUNT(*)
0	Overall its just OK when considering the price...	2
1	My wife has recurring extreme muscle spasms, u...	3
2	This coffee is horrible and unfortunately not ...	2
3	This will be the bottle that you grab from the...	3
4	I didnt like this coffee. Instead of telling y...	2

```
In [5]: display[display['UserId']=='AZY10LLTJ71NX']
```

```
Out [5]:
```

	UserId	ProductId	ProfileName	Time	\
80638	AZY10LLTJ71NX	B006P7E5ZI	undertheshrine "undertheshrine"	1334707200	

	Score	Text	COUNT(*)
80638	5	I was recommended to try green tea extract to ...	5

```
In [6]: display['COUNT(*)'].sum()
```

```
Out [6]: 393063
```

3 [2] Exploratory Data Analysis

3.1 [2.1] Data Cleaning: Deduplication

It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

```
In [7]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND UserId="AR5J8UI46CURR"
ORDER BY ProductID
""", con)
display.head()
```

```
Out [7]:
```

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	\
0	78445	B000HDL1RQ	AR5J8UI46CURR	Geetha Krishnan	2	
1	138317	B000HDOPYC	AR5J8UI46CURR	Geetha Krishnan	2	
2	138277	B000HDOPYM	AR5J8UI46CURR	Geetha Krishnan	2	

3	73791	B000HDOPZG	AR5J8UI46CURR	Geetha Krishnan	2
4	155049	B000PAQ75C	AR5J8UI46CURR	Geetha Krishnan	2

	HelpfulnessDenominator	Score	Time \
0	2	5	1199577600
1	2	5	1199577600
2	2	5	1199577600
3	2	5	1199577600
4	2	5	1199577600

	Summary \
0	LOACKER QUADRATINI VANILLA WAFERS
1	LOACKER QUADRATINI VANILLA WAFERS
2	LOACKER QUADRATINI VANILLA WAFERS
3	LOACKER QUADRATINI VANILLA WAFERS
4	LOACKER QUADRATINI VANILLA WAFERS

	Text
0	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
1	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
2	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
3	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
4	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...

As it can be seen above that same user has multiple reviews with same values for HelpfulnessNumerator, HelpfulnessDenominator, Score, Time, Summary and Text and on doing analysis it was found that ProductId=B000HDOPZG was Loacker Quadratini Vanilla Wafer Cookies, 8.82-Ounce Packages (Pack of 8) ProductId=B000HDL1RQ was Loacker Quadratini Lemon Wafer Cookies, 8.82-Ounce Packages (Pack of 8) and so on

It was inferred after analysis that reviews with same parameters other than ProductId belonged to the same product just having different flavour or quantity. Hence in order to reduce redundancy it was decided to eliminate the rows having same parameters.

The method used for the same was that we first sort the data according to ProductId and then just keep the first similar product review and delete the others. for eg. in the above just the review for ProductId=B000HDL1RQ remains. This method ensures that there is only one representative for each product and deduplication without sorting would lead to possibility of different representatives still existing for the same product.

```
In [8]: #Sorting data according to ProductId in ascending order
sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True, inplace=False)
```

```
In [9]: #Deduplication of entries
final=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"}, keep='first')
final.shape
```

```
Out[9]: (348262, 10)
```

```
In [10]: #Checking to see how much % of data still remains
(final['Id'].size*1.0)/(filtered_data['Id'].size*1.0)*100
```

Out[10]: 69.6524

Observation:- It was also seen that in two rows given below the value of HelpfulnessNumerator is greater than HelpfulnessDenominator which is not practically possible hence these two rows too are removed from calculations

```
In [11]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND Id=44737 OR Id=64422
ORDER BY ProductID
""", con)
```

```
display.head()
```

```
Out[11]:
```

	Id	ProductId	UserId	ProfileName	\
0	64422	B000MIDROQ	A161DK06JJMCYF	J. E. Stephens	"Jeanne"
1	44737	B001EQ55RW	A2V0I904FH7ABY		Ram

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	\
0		3	1	5	1224892800
1		3	2	4	1212883200

	Summary	\
0	Bought This for My Son at College	
1	Pure cocoa taste with crunchy almonds inside	

	Text
0	My son loves spaghetti so I didn't hesitate or...
1	It was almost a 'love at first bite' - the per...

```
In [12]: final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]
```

```
In [13]: #Before starting the next phase of preprocessing lets see the number of entries left
print(final.shape)
```

```
#How many positive and negative reviews are present in our dataset?
```

```
final['Score'].value_counts()
```

```
(348260, 10)
```

```
Out[13]: 1    293516
0     54744
Name: Score, dtype: int64
```

4 [3] Preprocessing

4.1 [3.1]. Preprocessing Review Text

Now that we have finished deduplication our data requires some preprocessing before we go on further with analysis and making the prediction model.

Hence in the Preprocessing phase we do the following in the order below:-

1. Begin by removing the html tags
2. Remove any punctuations or limited set of special characters like , or . or # etc.
3. Check if the word is made up of english letters and is not alpha-numeric
4. Check to see if the length of the word is greater than 2 (as it was researched that there is no adjective in 2-letters)
5. Convert the word to lowercase
6. Remove Stopwords
7. Finally Snowball Stemming the word (it was observed to be better than Porter Stemming)

After which we collect the words used to describe positive and negative reviews

```
In [14]: # printing some random reviews
sent_0 = final['Text'].values[0]
print(sent_0)
print("="*50)

sent_1000 = final['Text'].values[1000]
print(sent_1000)
print("="*50)

sent_1500 = final['Text'].values[1500]
print(sent_1500)
print("="*50)

sent_4900 = final['Text'].values[4900]
print(sent_4900)
print("="*50)
```

This book was purchased as a birthday gift for a 4 year old boy. He squealed with delight and I

=====

I've purchased both the Espresso Espresso (classic) and the 100% Arabica. My vote is defin

=====

This is a great product. It is very healthy for all of our dogs, and it is the first food that

=====

I find everything I need at Amazon so I always look there first. Chocolate tennis balls for a t

=====

```
In [15]: # remove urls from text python: https://stackoverflow.com/a/40823105/4084039
sent_0 = re.sub(r"http\S+", "", sent_0)
sent_1000 = re.sub(r"http\S+", "", sent_1000)
```

```

sent_150 = re.sub(r"http\S+", "", sent_1500)
sent_4900 = re.sub(r"http\S+", "", sent_4900)

print(sent_0)

```

This book was purchased as a birthday gift for a 4 year old boy. He squealed with delight and

```

In [16]: # https://stackoverflow.com/questions/16206380/python-beautifulsoup-how-to-remove-all
from bs4 import BeautifulSoup

soup = BeautifulSoup(sent_0, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1000, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1500, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_4900, 'lxml')
text = soup.get_text()
print(text)

```

This book was purchased as a birthday gift for a 4 year old boy. He squealed with delight and

I've purchased both the Espressione Espresso (classic) and the 100% Arabica. My vote is defin

This is a great product. It is very healthy for all of our dogs, and it is the first food that

I find everything I need at Amazon so I always look there first. Chocolate tennis balls for a

```

In [17]: # https://stackoverflow.com/a/47091490/4084039
import re

def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"n't", " not", phrase)

```



```

phrase = re.sub(r"\ 're", " are", phrase)
phrase = re.sub(r"\ 's", " is", phrase)
phrase = re.sub(r"\ 'd", " would", phrase)
phrase = re.sub(r"\ 'll", " will", phrase)
phrase = re.sub(r"\ 't", " not", phrase)
phrase = re.sub(r"\ 've", " have", phrase)
phrase = re.sub(r"\ 'm", " am", phrase)
return phrase

```

```

In [18]: sent_1500 = decontracted(sent_1500)
print(sent_1500)
print("="*50)

```

This is a great product. It is very healthy for all of our dogs, and it is the first food that
=====

```

In [19]: #remove words with numbers python: https://stackoverflow.com/a/18082370/4084039
sent_0 = re.sub("\S*\d\S*", "", sent_0).strip()
print(sent_0)

```

This book was purchased as a birthday gift for a year old boy. He squealed with delight and h

```

In [20]: #remove spacial character: https://stackoverflow.com/a/5843547/4084039
sent_1500 = re.sub('[^A-Za-z0-9]+', ' ', sent_1500)
print(sent_1500)

```

This is a great product It is very healthy for all of our dogs and it is the first food that t

```

In [21]: # https://gist.github.com/sebleier/554280
# we are removing the words from the stop words list: 'no', 'nor', 'not'
# <br /><br /> ==> after the above steps, we are getting "br br"
# we are including them into stop words list
# instead of <br /> if we have <br/> these tags would have reumoved in the 1st step

stopwords= set(['br', 'the', 'i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves',
                "you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', 'him',
                'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself',
                'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', 't',
                'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'h',
                'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'as',
                'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'through',
                'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off', 'o',
                'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'an',
                'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than', 'too',
                's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've", 'n',
                've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't",

```

```
"hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'mi',
"mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't",
'won', "won't", 'wouldn', "wouldn't"])
```

```
In [22]: # Combining all the above students
from tqdm import tqdm
preprocessed_reviews = []
# tqdm is for printing the status bar
for sentence in tqdm(final['Text'].values):
    sentence = re.sub(r"http\S+", "", sentence)
    sentence = BeautifulSoup(sentence, 'lxml').get_text()
    sentence = decontracted(sentence)
    sentence = re.sub("\S*\d\S*", "", sentence).strip()
    sentence = re.sub('[^A-Za-z]+', ' ', sentence)
    # https://gist.github.com/sebleier/554280
    sentence = ' '.join(e.lower() for e in sentence.split() if e.lower() not in stopwords)
    preprocessed_reviews.append(sentence.strip())
```

100%|| 348260/348260 [03:15<00:00, 1777.94it/s]

```
In [23]: preprocessed_reviews[1500]
```

```
Out[23]: 'great product healthy dogs first food love eat helped older dog lose weight year old'
```

```
In [24]: final['Text'] = preprocessed_reviews
```

```
#balancing the data
finalp = final[final.Score == 1].sample(50000,random_state =2)
finaln = final[final.Score == 0].sample(50000,random_state =2)
finalx = pd.concat([finalp,finaln],ignore_index=True)
finalx = finalx.sort_values('Time')
X = finalx.Text
```

5 [4] Featurization

5.1 [4.3] TF-IDF

```
In [25]: tf_idf_vect = TfidfVectorizer(min_df=10)
tfidf = tf_idf_vect.fit_transform(X)
print("some sample features(unique words in the corpus)",tf_idf_vect.get_feature_names())
print('='*50)
print("the type of count vectorizer ",type(tfidf))
print("the shape of out text TFIDF vectorizer ",tfidf.get_shape())
print("the number of unique words including both unigrams and bigrams ", tfidf.get_shape()[0])
```

```
some sample features(unique words in the corpus) ['aa', 'aafco', 'aback', 'abandon', 'abandoned',
=====
```

```
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
```

the shape of out text TFIDF vectorizer (100000, 12874)
the number of unique words including both unigrams and bigrams 12874

6 [5] Assignment 11: Truncated SVD

Apply Truncated-SVD on only this feature set:

SET 2:Review text, preprocessed one converted into vectors

Procedure:

Take top 2000 or 3000 features from tf-idf vectorizers using idf_ score.

You need to calculate the co-occurrence matrix with the selected features (Note: $X.X^T$

doesn't give the co-occurrence matrix, it returns the covariance matrix, check these blogs
1, blog-2 for more information)

You should choose the n_components in truncated svd, with maximum explained

variance. Please search on how to choose that and implement them. (hint: plot of cumulative
explained variance ratio)

After you are done with the truncated svd, you can apply K-Means clustering and choose

the best number of clusters based on elbow method.

 Print out wordclouds for each cluster, similar to that in previous assignment.

You need to write a function that takes a word and returns the most similar words using

cosine similarity between the vectors(vector: a row in the matrix after truncatedSVD)

6.1 Truncated-SVD

6.1.1 [5.1] Taking top features from TFIDF, SET 2

```
In [26]: indices = np.argsort(tf_idf_vect.idf_[:, -1])
```

```
In [27]: features = tf_idf_vect.get_feature_names()
```

```
In [28]: top_feat = [features[i] for i in indices[:2500]]
```

6.1.2 [5.2] Calculation of Co-occurrence matrix

```
In [29]: matrix = np.zeros((2500,2500))
```

```
In [30]: for i in tqdm(X):
    each_sentence = i.split()
    for index, words in enumerate(each_sentence):
        if words in top_feat:
            for x in range(max(index-5,0),min(index+5,len(each_sentence)-1)+1):
                if each_sentence[x] in top_feat:
                    matrix[top_feat.index(words),top_feat.index(each_sentence[x])] += 1
                else:
                    continue
        else:
            continue
```

```
100%|| 100000/100000 [10:04<00:00, 165.56it/s]
```


6.1.3 [5.3] Finding optimal value for number of components (n) to be retained.

```
In [540]: from sklearn.decomposition import TruncatedSVD
    model = TruncatedSVD(n_components=1250, n_iter=7, random_state=42)
    svd = model.fit_transform(matrix)
```

```
In [541]: evr = model.explained_variance_ratio_
```

```
In [542]: cum_evr = np.cumsum(evr)
```

```
In [543]: plt.plot(cum_evr)
    plt.grid()
    plt.xlabel('Number of Components')
    plt.ylabel('cumulative explained variance ratio')
    plt.show()
```



We observe that 100 components can explain 99% of the variance.

6.1.4 [5.4] Applying k-means clustering


```
In [544]: # Using 100 components as we got from S
          model = TruncatedSVD(n_components=100, n_iter=7, random_state=42)
          svd_final = model.fit_transform(matrix)

In [545]: svd_final.shape
Out[545]: (2500, 100)

In [546]: #We define a function for hyperparameter tuning of KMeans
          def kalgo(vector):
              ssd = []
              k = []
              for x in range(2,25,3):
                  kmeans = KMeans(n_clusters=x, random_state=0, n_jobs=-1).fit(vector)
                  ssd.append(kmeans.inertia_)
                  k.append(x)

          plt.figure()
          plt.plot(k,ssd,'bx-')
          plt.title('knee plot')
          plt.xlabel("Sum of squared distances")
          plt.ylabel("No. of clusters")
          plt.show()
```

```
In [400]: from sklearn.cluster import KMeans
          kalgo(svd_final)
```



As we can see from knee plot, $k=7$ is optimal k .

```
In [401]: model = KMeans(n_clusters=7, random_state=0, n_jobs=-1).fit(svd_final)
          w = model.labels_
```


```
In [402]: w.shape
```

```
Out[402]: (2500,)
```

6.1.5 [5.5] Wordclouds of clusters obtained in the above section

```
In [492]: #Create list of list in a variable name cluster of size max(w)+1. Append features to
          cluster = [[] for _ in range(max(w) + 1)]
          for i in range(0,2500):
              cluster[w[i]].append(top_feat[i])
          for i in range(0,max(w)+1):
              print("Wordcloud For cluster {}".format(i+1))
              wordcloud = WordCloud(background_color="white").generate(str(cluster[i]))
              plt.imshow(wordcloud)
              plt.axis("off")
              plt.show()
```

Wordcloud For cluster 1



Wordcloud For cluster 2

product'

Wordcloud For cluster 3

one'

Wordcloud For cluster 4

time'

Wordcloud For cluster 5

get'

Wordcloud For cluster 6

love'

Wordcloud For cluster 7

use'

6.1.6 [5.6] Function that returns most similar words for a given word.

```
In [549]: def cosine_sim(word):
          cmatrix = cosine_similarity(svd_final)
          similar_vect = cmatrix[top_feat.index(word)]
          print("Top 20 Words in order similar to '{}' are:".format(word))
          indices = similar_vect.argsort()[::-1][:20]
          #indices = np.argsort(similar_vect[::-1])[:20]
          for i in range(len(indices)):
              print(top_feat[indices[i]])

In [550]: #using word from above wordcloud
          cosine_sim('love')
```

Top 20 Words in order similar to 'love' are:

love
em
chihuahuas
grandkids
pugs
bueno
doggies
pina
kiddos
snobs
delish
nom
wedge
hazlenut
flatbread

sincere
boylan
yorkshire
slathered
uh

```
In [539]: #using from top_feat  
          cosine_sim(top_feat[1762])
```

Top 20 Words in order similar to 'product' are:

product
manner
page
marketed
disclaimer
meaningless
quoted
caliber
prompt
unreturnable
sketchy
specific
notification
cholesterol
marketing
consumers
claims
mechanical
development
advertise

7 [6] Conclusions

Optimal value for number of components : 100 Optimal value for number of clusters : 7

Conclusions:

- 1) Observed that SVD can manage with sparse matrix unlike PCA.
- 2) Using only the top components, we can cut down the cost of computation.
- 3) I have observed the cosine similarity with both the occurrence matrix and svd_final matrix. It means SVD performs very well in getting the top components.