

## assign3(2)

December 29, 2018

```
In [2]: import warnings
warnings.filterwarnings("ignore")
import sqlite3
import numpy as np
import pandas as pd
import nltk
import string
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score
from sklearn.model_selection import cross_val_score
from collections import Counter
from sklearn.metrics import accuracy_score
from sklearn import model_selection
from sklearn.metrics import f1_score
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer
from sklearn.metrics import classification_report, confusion_matrix
import re
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle
from tqdm import tqdm
import os
!pip install -q scikit-plot
import scikitplot.metrics as skplt
```

```
import sqlite3
from google.colab import drive
drive.mount('/content/drive/')
```

Drive already mounted at /content/drive/; to attempt to forcibly remount, call drive.mount("/c

```
In [0]: os.chdir("/content/drive/My Drive/Colab Notebooks")
```

```
In [0]: con = sqlite3.connect("final.sqlite")
```

```
In [0]: import pandas as pd
        filtered_data = pd.read_sql_query("""
        SELECT * FROM REVIEWS
        """, con)
```

```
In [6]: filtered_data.shape
```

```
Out[6]: (364171, 12)
```

```
In [0]: filtered_data.set_index = filtered_data.index
```

```
In [0]: filtered_data.drop(['index'],axis=1)
        #dropping index value from columns
```

```
In [0]: final_data = filtered_data.sample(60000,random_state=2)

        #sampling 100k datapoints
```

```
In [0]: final_data.head(3)
        final_data = final_data.sort_values('Time')
```

```
In [11]: final_data.Score.value_counts()
```

```
Out[11]: positive    50677
         negative     9323
         Name: Score, dtype: int64
```

```
In [0]: final_data.head(3)
```

```
In [0]: final_data_kd_tree = final_data.sample(20000,random_state=2)
        #sampling 20 k points for kd_tree
```

```
In [14]: final_data_kd_tree.Score.value_counts()
```

```
Out[14]: positive    16882
         negative     3118
         Name: Score, dtype: int64
```

## 1 KNN on AVG W2V using KD\_tree

```
In [0]: #training own model
```

```
i=0
list_of_sent=[]
for sent in filtered_data['CleanedText'].values:
    list_of_sent.append(sent.split())
```

```
In [0]: w2v_model=Word2Vec(list_of_sent,min_count=5,size=50,workers=4)
```

```
In [0]: w2v_words = list(w2v_model.wv.vocab)
```

```
In [18]: w2v_model.wv.most_similar('tasti')
```

```
Out[18]: [('delici', 0.8092088103294373),
          ('yummi', 0.7842735052108765),
          ('tastey', 0.7792528867721558),
          ('hearti', 0.6838183403015137),
          ('nutriti', 0.6770579218864441),
          ('satisfi', 0.6765849590301514),
          ('good', 0.6748582720756531),
          ('nice', 0.6626357436180115),
          ('terrif', 0.6532567143440247),
          ('crunchi', 0.6260974407196045)]
```

```
In [0]: X = final_data_kd_tree['CleanedText'].values
y = final_data_kd_tree.Score.values
X_tr, X_test, y_tr, y_test = model_selection.train_test_split(X, y, test_size=0.3, random_state=42)
# split the train data set into cross validation train and cross validation test
```

```
In [21]: #training data
```

```
i=0
list_of_sent=[]
for sent in X_tr:
    list_of_sent.append(sent.split())

#Avgw2v
sent_vectors = []
for sent in tqdm(list_of_sent):
    sent_vec = np.zeros(50)
    cnt_words=0
    for word in sent:
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    sent_vectors.append(sent_vec)
```

100%|| 14000/14000 [00:29<00:00, 478.08it/s]

```
In [33]: X_tr = sent_vectors
```

```
print(len(list_of_sent),len(X_tr))
```

14000 14000

```
In [34]: #test data
```

```
#Avgw2v
```

```
i=0
```

```
list_of_sent=[]
```

```
for sent in X_test:
```

```
    list_of_sent.append(sent.split())
```

```
#Avgw2v
```

```
sent_vectors = []
```

```
for sent in tqdm(list_of_sent):
```

```
    sent_vec = np.zeros(50)
```

```
    cnt_words=0
```

```
    for word in sent:
```

```
        if word in w2v_words:
```

```
            vec = w2v_model.wv[word]
```

```
            sent_vec += vec
```

```
            cnt_words += 1
```

```
    if cnt_words != 0:
```

```
        sent_vec /= cnt_words
```

```
    sent_vectors.append(sent_vec)
```

100%|| 6000/6000 [00:12<00:00, 468.24it/s]

```
In [36]: X_test = sent_vectors
```

```
print(len(list_of_sent),len(X_test))
```

6000 6000

```
In [37]: print(len(X_tr[0]))
```

50

```
In [39]: #Finding optimal k
```

```
# creating odd list of K for KNN
```

```
myList = list(range(5,30))
```

```

neighbors = list(filter(lambda x: x % 2 != 0, myList))

# empty list that will hold cv scores
cv_scores = []

# perform 2-fold cross validation
for k in neighbors:
    knn = KNeighborsClassifier(n_neighbors=k,algorithm='kd_tree')
    scores = cross_val_score(knn, X_tr, y_tr, cv=2, scoring='f1_micro')
    cv_scores.append(scores.mean())

# changing to misclassification error
MSE = [1 - x for x in cv_scores]

# determining best k
optimal_k = neighbors[MSE.index(min(MSE))]
print('\nThe optimal number of neighbors is %d.' % optimal_k)

# plot misclassification error vs k
plt.plot(neighbors, MSE)


for xy in zip(neighbors, np.round(MSE,3)):
    plt.annotate('(%s, %s)' % xy, xy=xy, textcoords='data')

plt.xlabel('Number of Neighbors K')
plt.ylabel('Misclassification Error')
plt.show()

print("the misclassification error for each k value is : ", np.round(MSE,3))

```

The optimal number of neighbors is 9.




the misclassification error for each k value is : [0.131 0.13 0.128 0.128 0.13 0.129 0.131 0.131 0.131]

```
In [40]: #Performing KNN using kd_tree
knn = KNeighborsClassifier(n_neighbors=optimal_k,algorithm='kd_tree')
knn.fit(X_tr,y_tr)
pred = knn.predict(X_test)
acc = f1_score(y_test, pred,average='micro') * float(100)
print('\nThe accuracy of the knn classifier for k = %d is %f%%' % (optimal_k, acc))
```

The accuracy of the knn classifier for k = 9 is 87.183333%

```
In [41]: skplt.plot_confusion_matrix(y_test ,pred)
```

```
Out[41]: <matplotlib.axes._subplots.AxesSubplot at 0x7f0507d6c748>
```



```
In [42]: from sklearn.metrics import classification_report
         print(classification_report(y_test ,pred))
```

	precision	recall	f1-score	support
negative	0.72	0.31	0.43	946
positive	0.88	0.98	0.93	5054
micro avg	0.87	0.87	0.87	6000
macro avg	0.80	0.64	0.68	6000
weighted avg	0.86	0.87	0.85	6000

## 2 KNN on AVG W2V using Brute Force

```
In [0]: X = final_data['CleanedText'].values
         y = final_data.Score.values
```

```
X_tr, X_test, y_tr, y_test = model_selection.train_test_split(X, y, test_size=0.3, random_state=42)
# split the train data set into cross validation train and cross validation test
```

```
In [44]: #training data
i=0
list_of_sent=[]
for sent in X_tr:
    list_of_sent.append(sent.split())

#Avgw2v
sent_vectors = []
for sent in tqdm(list_of_sent):
    sent_vec = np.zeros(50)
    cnt_words=0
    for word in sent:
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    sent_vectors.append(sent_vec)
```

```
100%| 42000/42000 [01:28<00:00, 476.39it/s]
```

```
In [45]: X_tr = sent_vectors

print(len(list_of_sent),len(X_tr))
```

```
42000 42000
```

```
In [46]: #test data
#Avgw2v
i=0
list_of_sent=[]
for sent in X_test:
    list_of_sent.append(sent.split())

#Avgw2v
sent_vectors = []
for sent in tqdm(list_of_sent):
    sent_vec = np.zeros(50)
    cnt_words=0
    for word in sent:
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
```



```

        cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    sent_vectors.append(sent_vec)

```

100%|| 18000/18000 [00:37<00:00, 482.83it/s]

In [47]: X\_test = sent\_vectors

```

    print(len(list_of_sent),len(X_test))

```

18000 18000

In [48]: print(len(X\_tr[0]))

50

```

In [49]: #finding optimal k
         # creating odd list of K for KNN
myList = list(range(5,30))
neighbors = list(filter(lambda x: x % 2 != 0, myList))

# empty list that will hold cv scores
cv_scores = []

# perform 2-fold cross validation
for k in neighbors:
    knn = KNeighborsClassifier(n_neighbors=k)
    scores = cross_val_score(knn, X_tr, y_tr, cv=2, scoring='f1_micro')
    cv_scores.append(scores.mean())

# changing to misclassification error
MSE = [1 - x for x in cv_scores]

# determining best k
optimal_k = neighbors[MSE.index(min(MSE))]
print('\nThe optimal number of neighbors is %d.' % optimal_k)

# plot misclassification error vs k
plt.plot(neighbors, MSE)

for xy in zip(neighbors, np.round(MSE,3)):
    plt.annotate('%s, %s' % xy, xy=xy, textcoords='data')


plt.xlabel('Number of Neighbors K')
plt.ylabel('Misclassification Error')

```

```
plt.show()
```

```
print("the misclassification error for each k value is : ", np.round(MSE,3))
```

The optimal number of neighbors is 11.




the misclassification error for each k value is : [0.124 0.121 0.121 0.12 0.121 0.121 0.121 0.124]

```
In [50]: #Performing KNN using brute force
knn = KNeighborsClassifier(n_neighbors=optimal_k)
knn.fit(X_tr,y_tr)
pred = knn.predict(X_test)
acc = f1_score(y_test, pred,average='micro') * float(100)
print('\nThe accuracy of the knn classifier for k = %d is %f%%' % (optimal_k, acc))
```

The accuracy of the knn classifier for k = 11 is 88.022222%

```
In [51]: skplt.plot_confusion_matrix(y_test ,pred)
```

Out [51]: <matplotlib.axes.\_subplots.AxesSubplot at 0x7f04e1c98828>



```
In [52]: print(classification_report(y_test ,pred))
```

	precision	recall	f1-score	support
negative	0.78	0.34	0.47	2851
positive	0.89	0.98	0.93	15149
micro avg	0.88	0.88	0.88	18000
macro avg	0.83	0.66	0.70	18000
weighted avg	0.87	0.88	0.86	18000

### 3 KNN on weighted TFIDF using KD\_tree

```
In [0]: X = final_data_kd_tree['CleanedText'].values  
        y = final_data_kd_tree.Score.values
```

```
X_tr, X_test, y_tr, y_test = model_selection.train_test_split(X, y, test_size=0.3, random_state=42)
# split the train data set into cross validation train and cross validation test
```

```
In [58]: i=0
list_of_sent=[]
for sent in X_tr:
    list_of_sent.append(sent.split())

i=0
list_of_sent_test=[]
for sent in X_test:
    list_of_sent_test.append(sent.split())

len(list_of_sent)
```

```
Out[58]: 14000
```

```
In [0]: model = TfidfVectorizer()
tf_idf_matrix = model.fit_transform(X_tr)
tf2 = model.transform(X_test)
dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))
```

```
In [0]: # TF-IDF weighted Word2Vec
tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in tqdm(list_of_sent): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            # tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole corpus
            # sent.count(word) = tf value of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_sent_vectors.append(sent_vec)
    row += 1
```

```
In [61]: X_tr =tfidf_sent_vectors
len(X_tr)
```

Out [61]: 14000

```
In [0]: tfidf_feat = model.get_feature_names() # tfidf words/col-names
        # final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

        tfidf_sent_vectors_test = []; # the tfidf-w2v for each sentence/review is stored in th
        row=0;
        for sent in list_of_sent_test: # for each review/sentence
            sent_vec = np.zeros(50) # as word vectors are of zero length
            weight_sum = 0; # num of words with a valid vector in the sentence/review
            for word in sent: # for each word in a review/sentence
                try:
                    vec = w2v_model.wv[word]
                    # obtain the tf_idfidf of a word in a sentence/review
                    tfidf = final_tf_idf[row, tfidf_feat.index(word)]
                    sent_vec += (vec * tfidf)
                    weight_sum += tfidf
                except:
                    pass
            sent_vec /= weight_sum
            tfidf_sent_vectors_test.append(sent_vec)
            row += 1
```

```
In [63]: X_test = tfidf_sent_vectors_test
        len(X_test)
```

Out [63]: 6000

```
In [0]: X_tr = np.nan_to_num(X_tr)
        X_test = np.nan_to_num(X_test)
```

```
In [65]: myList = list(range(5,30))
        neighbors = list(filter(lambda x: x % 2 != 0, myList))

        # empty list that will hold cv scores
        cv_scores = []

        # perform 2-fold cross validation
        for k in neighbors:
            knn = KNeighborsClassifier(n_neighbors=k, algorithm='kd_tree')
            scores = cross_val_score(knn, X_tr, y_tr, cv=2, scoring='f1_micro')
            cv_scores.append(scores.mean())

        # changing to misclassification error
        MSE = [1 - x for x in cv_scores]

        # determining best k
        optimal_k = neighbors[MSE.index(min(MSE))]
        print('\nThe optimal number of neighbors is %d.' % optimal_k)
```

```

# plot misclassification error vs k
plt.plot(neighbors, MSE)

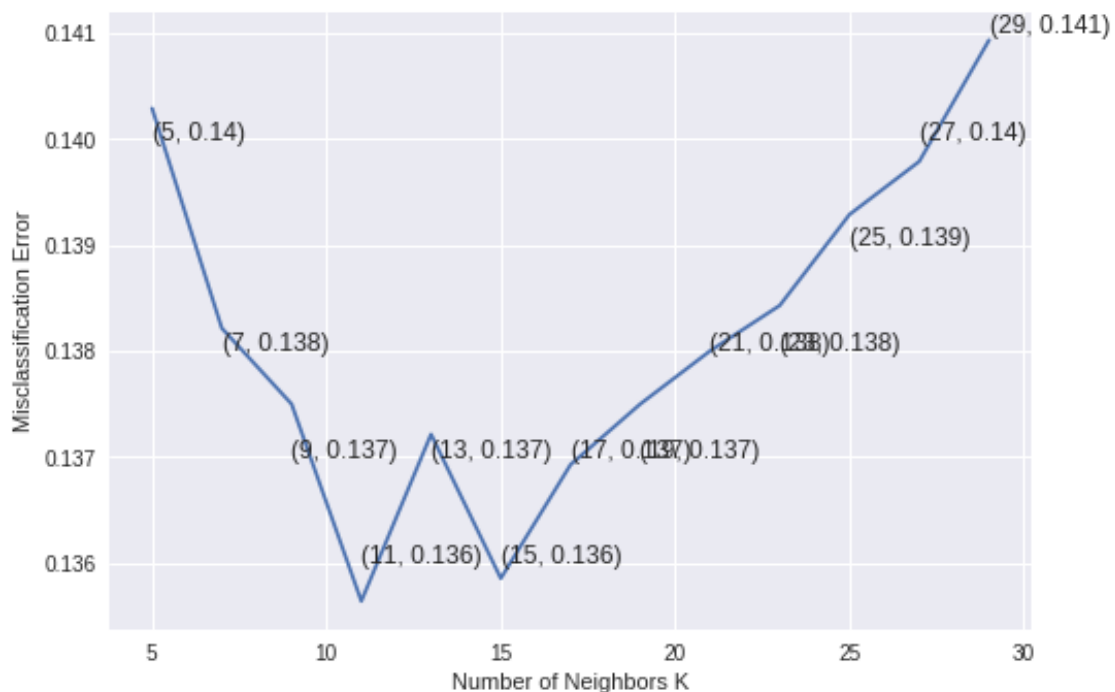
for xy in zip(neighbors, np.round(MSE,3)):
    plt.annotate('%s, %s' % xy, xy=xy, textcoords='data')

plt.xlabel('Number of Neighbors K')
plt.ylabel('Misclassification Error')
plt.show()

print("the misclassification error for each k value is : ", np.round(MSE,3))

```

The optimal number of neighbors is 11.



the misclassification error for each k value is : [0.14 0.138 0.137 0.136 0.137 0.136 0.137 0.141]

```

In [66]: #Performing KNN using kd-tree
knn = KNeighborsClassifier(n_neighbors=optimal_k,algorithm='kd_tree')
knn.fit(X_tr,y_tr)
pred = knn.predict(X_test)


```

```
acc = f1_score(y_test, pred,average='micro') * float(100)
print('\nThe accuracy of the knn classifier for k = %d is %f%%' % (optimal_k, acc))
```

The accuracy of the knn classifier for k = 11 is 84.066667%

```
In [67]: skplt.plot_confusion_matrix(y_test ,pred)
```

```
Out[67]: <matplotlib.axes._subplots.AxesSubplot at 0x7f95f777bfd0>
```



```
In [68]: print(classification_report(y_test ,pred))
```

	precision	recall	f1-score	support
negative	0.00	0.00	0.00	956
positive	0.84	1.00	0.91	5044
micro avg	0.84	0.84	0.84	6000

macro avg	0.42	0.50	0.46	6000
weighted avg	0.71	0.84	0.77	6000

## 4 KNN on weighted TFIDF using brute force

```
In [0]: X = final_data['CleanedText'].values
        y = final_data.Score.values
        X_tr, X_test, y_tr, y_test = model_selection.train_test_split(X, y, test_size=0.3, random_state=42)
        # split the train data set into cross validation train and cross validation test
```

```
In [70]: i=0
        list_of_sent=[]
        for sent in X_tr:
            list_of_sent.append(sent.split())

        i=0
        list_of_sent_test=[]
        for sent in X_test:
            list_of_sent_test.append(sent.split())

        len(list_of_sent)
```

```
Out[70]: 42000
```

```
In [0]: model = TfidfVectorizer()
        tf1 = model.fit_transform(X_tr)
        tf2 = model.transform(X_test)
        dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))
```

```
In [0]: # TF-IDF weighted Word2Vec
        tfidf_feat = model.get_feature_names() # tfidf words/col-names
        # final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

        tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in this list
        row=0;
        for sent in tqdm(list_of_sent): # for each review/sentence
            sent_vec = np.zeros(50) # as word vectors are of zero length
            weight_sum = 0; # num of words with a valid vector in the sentence/review
            for word in sent: # for each word in a review/sentence
                if word in w2v_words:
                    vec = w2v_model.wv[word]
                    # tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
                    # to reduce the computation we are
                    # dictionary[word] = idf value of word in whole corpus
                    # sent.count(word) = tf value of word in this review
                    tf_idf = dictionary[word]*(sent.count(word)/len(sent))
```



```

        sent_vec += (vec * tf_idf)
        weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_sent_vectors.append(sent_vec)
    row += 1

```

```

In [73]: X_tr = tfidf_sent_vectors
        len(X_tr)

```

```

Out[73]: 42000

```

```

In [0]: # TF-IDF weighted Word2Vec
tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in tqdm(list_of_sent_test): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        try:
            if word in w2v_words:
                vec = w2v_model.wv[word]
                tf_idf = dictionary[word]*(sent.count(word)/len(sent))
                sent_vec += (vec * tf_idf)
                weight_sum += tf_idf
        except:
            pass
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_sent_vectors.append(sent_vec)
    row += 1

```

```

In [85]: X_test = tfidf_sent_vectors
        len(X_test)

```

```

Out[85]: 18000

```

```

In [0]: X_tr = np.nan_to_num(X_tr)
        X_test = np.nan_to_num(X_test)

```

```

In [87]: myList = list(range(5,30))
        neighbors = list(filter(lambda x: x % 2 != 0, myList))

        # empty list that will hold cv scores
        cv_scores = []

```

```

# perform 2-fold cross validation
for k in neighbors:
    knn = KNeighborsClassifier(n_neighbors=k)
    scores = cross_val_score(knn, X_tr, y_tr, cv=2, scoring='f1_micro')
    cv_scores.append(scores.mean())

# changing to misclassification error
MSE = [1 - x for x in cv_scores]

# determining best k
optimal_k = neighbors[MSE.index(min(MSE))]
print('\nThe optimal number of neighbors is %d.' % optimal_k)

# plot misclassification error vs k
plt.plot(neighbors, MSE)

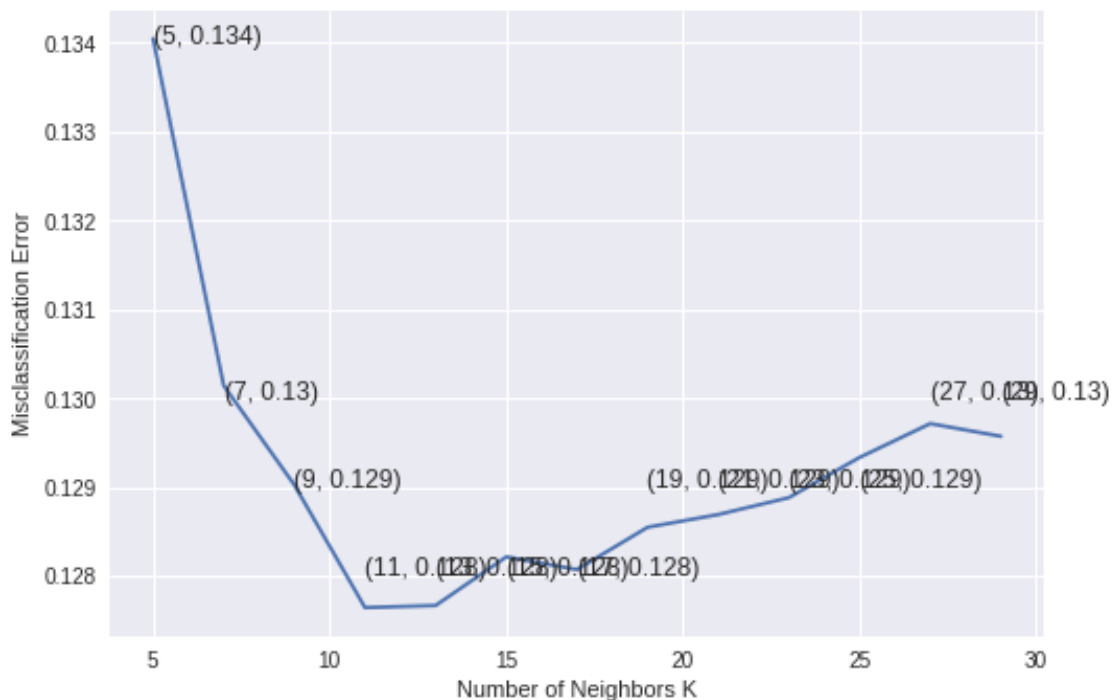
for xy in zip(neighbors, np.round(MSE,3)):
    plt.annotate('(%s, %s)' % xy, xy=xy, textcoords='data')

plt.xlabel('Number of Neighbors K')
plt.ylabel('Misclassification Error')
plt.show()

print("the misclassification error for each k value is : ", np.round(MSE,3))

```

The optimal number of neighbors is 11.




the misclassification error for each k value is : [0.134 0.13 0.129 0.128 0.128 0.128 0.128 0.128 0.13 ]

```
In [88]: #Performing KNN using brute force
knn = KNeighborsClassifier(n_neighbors=optimal_k)
knn.fit(X_tr,y_tr)
pred = knn.predict(X_test)
acc = f1_score(y_test, pred,average='micro') * float(100)
print('\nThe accuracy of the knn classifier for k = %d is %f%%' % (optimal_k, acc))
```

The accuracy of the knn classifier for k = 11 is 86.988889%

```
In [89]: skplt.plot_confusion_matrix(y_test ,pred)
```

```
Out[89]: <matplotlib.axes._subplots.AxesSubplot at 0x7f95f5a96d68>
```



```
In [90]: print(classification_report(y_test ,pred))
```

```

              precision    recall  f1-score   support

 negative     0.76      0.28      0.41      2912
 positive     0.88      0.98      0.93     15088

 micro avg     0.87      0.87      0.87     18000
 macro avg     0.82      0.63      0.67     18000
 weighted avg   0.86      0.87      0.84     18000

```

```
In [0]: !pip install -q PTable
```

```
In [0]: from prettytable import PrettyTable
```

```
In [0]: z = PrettyTable()
```

```
z.field_names = ["Vectorizer", "Model", "Hyperparameter k", "f1 score accuracy"]
```

```
In [95]: #Final summary
```

```

z.add_row(["BoW", 'kd_tree', 29, '91.65%'])
z.add_row(["BoW", 'brute_force', 11, '91.66%'])
z.add_row(["TF_IDF", 'kd_tree', 23, '91.87%'])
z.add_row(["TF_IDF", 'brute_force', 9, '92.25%'])
z.add_row(["Avg W2V", 'kd_tree', 9, '87.18%'])
z.add_row(["Avg W2V", 'brute_force', 11, '88.02%'])
z.add_row(["TF_IDF weighted W2V", 'kd_tree', 11, '84.06%'])
z.add_row(["TF_IDF weighted W2V", 'brute_force', 11, '86.98%'])
print(z)

```

Vectorizer	Model	Hyperparameter k	f1 score accuracy
BoW	kd_tree	29	91.65%
BoW	brute_force	11	91.66%
TF_IDF	kd_tree	23	91.87%
TF_IDF	brute_force	9	92.25%
Avg W2V	kd_tree	9	87.18%
Avg W2V	brute_force	11	88.02%
TF_IDF weighted W2V	kd_tree	11	84.06%
TF_IDF weighted W2V	brute_force	11	86.98%

Conclusions: Different size of data points were taken for kd\_tree and brute force. Hence, the conclusion could be biased a bit. Brute Force: 100k Kd\_tree: 20k 2 fold CV is used for both algorithms instead of simple CV

1)It takes more time to run kd\_tree than brute force even after reducing the dimensions. But considering this set of points of about 100k for brute force, even this takes more time.

2)Knn on TFIDF Weighted avgw2v using kd\_tree model is a dumb model,here. It actually predicted only the positive class,the majority class, even if it gives a good f1 accuracy of 84%.