

assign3

December 29, 2018

1 Assignment 3:

1.1 Objective :

1.1.1 Apply Brute force KNN and KDTree on BoW,TFIDF,AVGW2V and Weighted TFIDF.

1.1.2 Get the conclusion table.

```
In [1]: import warnings
        warnings.filterwarnings("ignore")
        import sqlite3
        import numpy as np
        import pandas as pd
        import nltk
        import string
        from sklearn.model_selection import train_test_split
        from sklearn.neighbors import KNeighborsClassifier
        from sklearn.metrics import accuracy_score
        from sklearn.model_selection import cross_val_score
        from collections import Counter
        from sklearn.metrics import accuracy_score
        from sklearn import model_selection
        from sklearn.metrics import f1_score
        import matplotlib.pyplot as plt
        import seaborn as sns
        from sklearn.feature_extraction.text import TfidfTransformer
        from sklearn.feature_extraction.text import TfidfVectorizer
        from sklearn.feature_extraction.text import CountVectorizer
        from sklearn.metrics import confusion_matrix
        from sklearn import metrics
        from sklearn.metrics import roc_curve, auc
        from nltk.stem.porter import PorterStemmer
        from sklearn.metrics import classification_report, confusion_matrix
        import re
        from nltk.corpus import stopwords
        from nltk.stem import PorterStemmer
        from nltk.stem.wordnet import WordNetLemmatizer
```

```

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle
from tqdm import tqdm
import os

import sqlite3
from google.colab import drive
drive.mount('/content/drive/')

```

Drive already mounted at /content/drive/; to attempt to forcibly remount, call drive.mount("/c

2 Data Preprocessing

```

In [0]: os.chdir("/content/drive/My Drive/Colab Notebooks")

In [0]: con = sqlite3.connect("final.sqlite")

In [0]: import pandas as pd
        filtered_data = pd.read_sql_query("""
        SELECT * FROM REVIEWS
        """, con)

In [5]: filtered_data.shape

Out[5]: (364171, 12)

In [0]: filtered_data.set_index = filtered_data.index

In [0]: filtered_data.drop(['index'],axis=1)
        #dropping index value from columns

In [0]: final_data = filtered_data.sample(100000,random_state=2)

        #sampling 100k datapoints

In [0]: final_data.head(3)
        final_data = final_data.sort_values('Time')

In [10]: final_data.Score.value_counts()

Out[10]: 1    84541
         0    15459
         Name: Score, dtype: int64

In [0]: final_data_kd_tree = final_data.sample(20000,random_state=2)
        #sampling 20 k points for kd_tree

In [15]: final_data_kd_tree.Score.value_counts()

Out[15]: 1    16917
         0     3083
         Name: Score, dtype: int64

```

3 KNN on BoW using KD_Tree

```
In [0]: count_vect = CountVectorizer(ngram_range=(1,2))
        X = final_data_kd_tree['CleanedText'].values

        X_tr, X_test = model_selection.train_test_split(X, test_size=0.3, random_state=0)
        # split the train data set into cross validation train and cross validation test
        #X_tr, X_cv= model_selection.train_test_split(X_1, test_size=0.3, random_state=0)
        final_counts = count_vect.fit_transform(X_tr)
```

```
In [24]: print("the shape of out text BOW vectorizer ",final_counts.shape)
        print("the number of unique words are ", final_counts.get_shape()[1])
```

```
the shape of out text BOW vectorizer (14000, 305221)
the number of unique words are 305221
```

```
In [25]: #using truncated svd for dimensionality reduction
```

```
        from sklearn.decomposition import TruncatedSVD
        svd = TruncatedSVD(n_components=300)
        svd_data=svd.fit_transform(final_counts)
        svd_data.shape
```

```
Out[25]: (14000, 300)
```

Taking 300 dimensions for KD_TREE. Tried to plot (Explained Variance, No. of features retained), but it gives error for numpy. Says no attribute for explained_variance_.

```
In [0]: y = final_data_kd_tree['Score'].values

        y_tr, y_test = model_selection.train_test_split(y, test_size=0.3, random_state=0)
```

```
In [0]: #for CV and test data
        final_counts_test = count_vect.transform(X_test)
        svd_data_test=svd.transform(final_counts_test)
```

```
In [30]: svd_data_test.shape
```

```
Out[30]: (6000, 300)
```

Performing 2 cross validation. Finding the optimal k as follows. Using 'f1' score for accuracy.

```
In [35]: # creating odd list of K for KNN
        myList = list(range(5,30))
        neighbors = list(filter(lambda x: x % 2 != 0, myList))

        # empty list that will hold cv scores
        cv_scores = []
```

```

# perform 2-fold cross validation
for k in neighbors:
    knn = KNeighborsClassifier(n_neighbors=k,algorithm='kd_tree')
    scores = cross_val_score(knn, svd_data, y_tr, cv=2, scoring='f1')
    cv_scores.append(scores.mean())

# changing to misclassification error
MSE = [1 - x for x in cv_scores]

# determining best k
optimal_k = neighbors[MSE.index(min(MSE))]
print('\nThe optimal number of neighbors is %d.' % optimal_k)

# plot misclassification error vs k
plt.plot(neighbors, MSE)

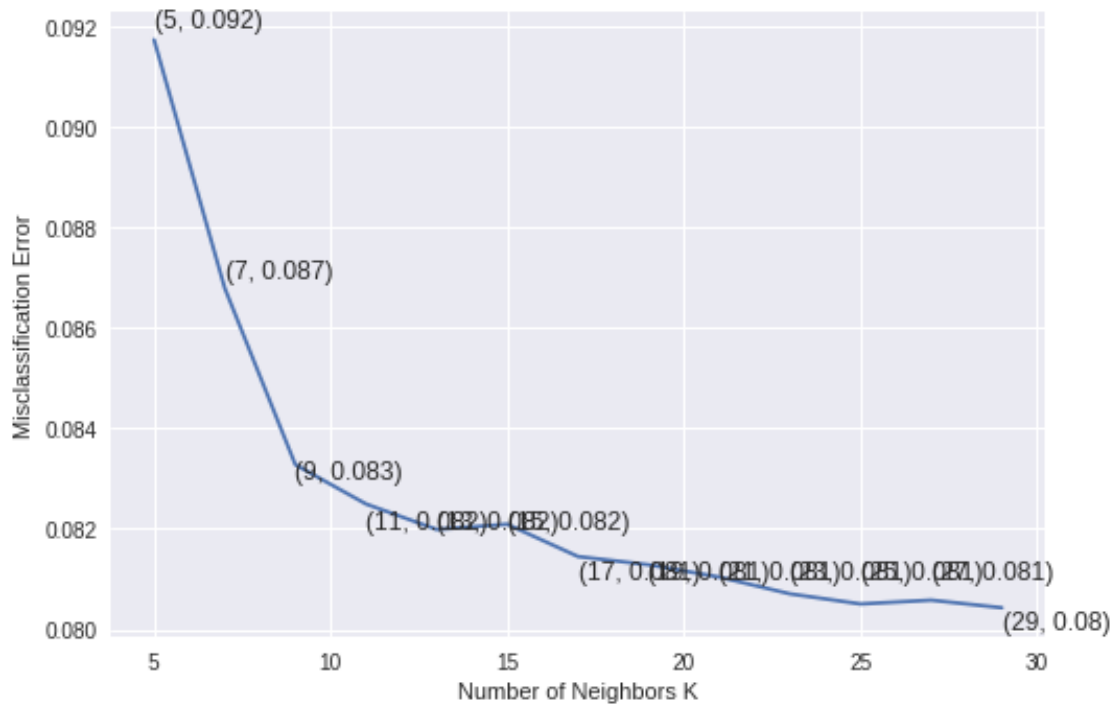
for xy in zip(neighbors, np.round(MSE,3)):
    plt.annotate('(%s, %s)' % xy, xy=xy, textcoords='data')

plt.xlabel('Number of Neighbors K')
plt.ylabel('Misclassification Error')
plt.show()

print("the misclassification error for each k value is : ", np.round(MSE,3))

```

The optimal number of neighbors is 29.



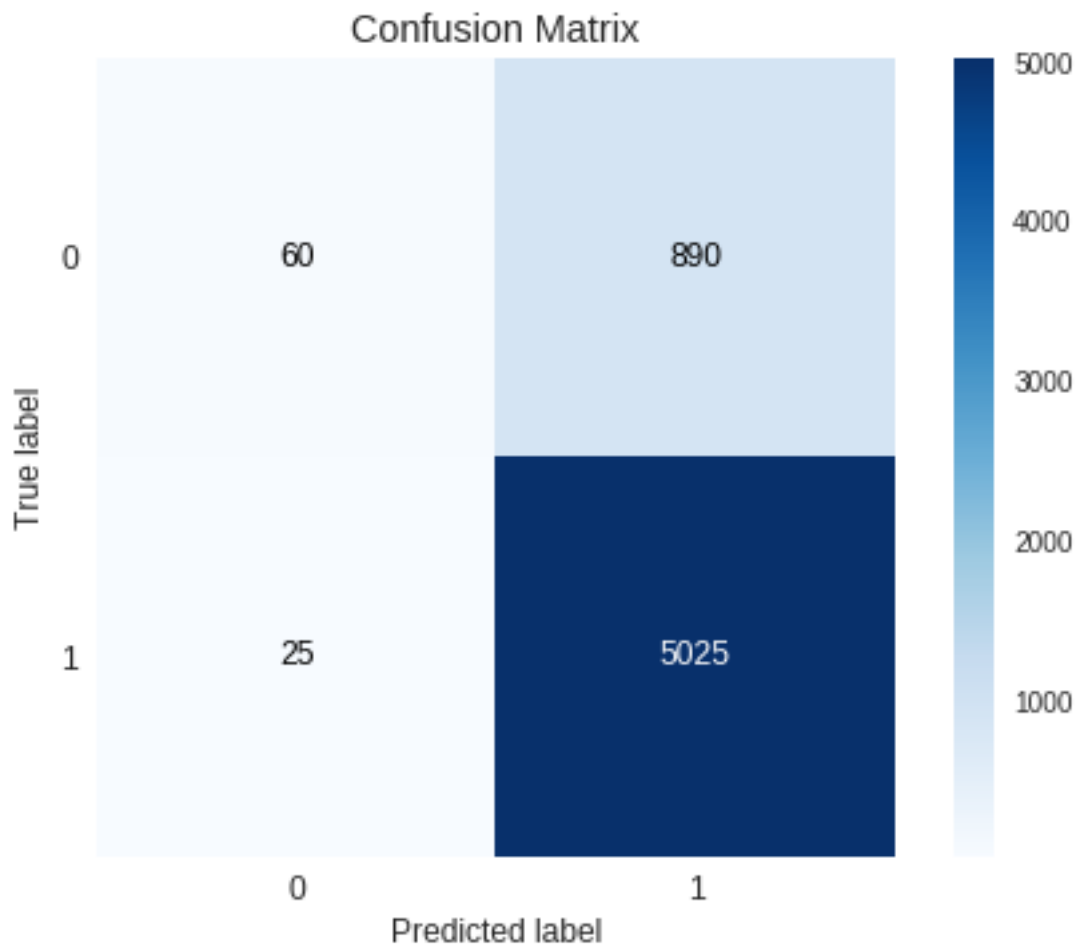
the misclassification error for each k value is : [0.092 0.087 0.083 0.082 0.082 0.082 0.081 0.08]

```
In [36]: #Performing KNN using kd_tree
knn = KNeighborsClassifier(n_neighbors=optimal_k,algorithm='kd_tree')
knn.fit(svd_data,y_tr)
pred = knn.predict(svd_data_test)
acc = f1_score(y_test, pred) * float(100)
print('\nThe accuracy of the knn classifier for k = %d is %f%%' % (optimal_k, acc))
```

The accuracy of the knn classifier for k = 29 is 91.655267%

```
In [51]: import scikitplot as skplt
skplt.metrics.plot_confusion_matrix(y_test ,pred)
```

```
Out[51]: <matplotlib.axes._subplots.AxesSubplot at 0x7f06e91f5e80>
```



```
In [52]: print(classification_report(y_test ,pred))
```

	precision	recall	f1-score	support
0	0.71	0.06	0.12	950
1	0.85	1.00	0.92	5050
micro avg	0.85	0.85	0.85	6000
macro avg	0.78	0.53	0.52	6000
weighted avg	0.83	0.85	0.79	6000

4 KNN using Brute Force on BoW

```
In [0]: #Performing KNN using Brute Force
        # split the data set into train and test
```

```

X = final_data['CleanedText'].values
y = final_data.Score.values
X_tr, X_test, y_tr, y_test = model_selection.train_test_split(X, y, test_size=0.3, random_state=42)
# split the train data set into cross validation train and cross validation test
count_vect = CountVectorizer(ngram_range=(1,2))
X_tr = count_vect.fit_transform(X_tr)
X_test = count_vect.transform(X_test)

```

In [58]: *#Getting the optimal k*

```

# creating odd list of K for KNN using Brute Force
myList = list(range(5,40))
neighbors = list(filter(lambda x: x % 2 != 0, myList))

# empty list that will hold cv scores
cv_scores = []

# perform 2-fold cross validation
for k in neighbors:
    knn = KNeighborsClassifier(n_neighbors=k)
    scores = cross_val_score(knn, X_tr, y_tr, cv=2, scoring='f1')
    cv_scores.append(scores.mean())

# changing to misclassification error
MSE = [1 - x for x in cv_scores]

# determining best k
optimal_k = neighbors[MSE.index(min(MSE))]
print('\nThe optimal number of neighbors is %d.' % optimal_k)

# plot misclassification error vs k
plt.plot(neighbors, MSE)


for xy in zip(neighbors, np.round(MSE,3)):
    plt.annotate('%s, %s' % xy, xy=xy, textcoords='data')

plt.xlabel('Number of Neighbors K')
plt.ylabel('Misclassification Error')
plt.show()

print("the misclassification error for each k value is : ", np.round(MSE,3))

```

The optimal number of neighbors is 11.




the misclassification error for each k value is : [0.084 0.084 0.083 0.083 0.083 0.083 0.084 0.084 0.084 0.084 0.084 0.084]

```
In [59]: #Performing KNN using brute force
knn = KNeighborsClassifier(n_neighbors=optimal_k)
knn.fit(X_tr,y_tr)
pred = knn.predict(X_test)
acc = f1_score(y_test, pred) * float(100)
print('\nThe accuracy of the knn classifier for k = %d is %f%%' % (optimal_k, acc))
```

The accuracy of the knn classifier for k = 11 is 91.663802%

```
In [61]: skplt.metrics.plot_confusion_matrix(y_test ,pred)
```

```
Out[61]: <matplotlib.axes._subplots.AxesSubplot at 0x7f06d66346d8>
```

```
In [62]: print(classification_report(y_test ,pred))
```

	precision	recall	f1-score	support
0	0.70	0.01	0.03	4644
1	0.85	1.00	0.92	25356
micro avg	0.85	0.85	0.85	30000
macro avg	0.78	0.51	0.47	30000
weighted avg	0.82	0.85	0.78	30000

5 Knn on TFIDF using KD_tree

```
In [63]: X = final_data_kd_tree['CleanedText'].values
         y = final_data_kd_tree.Score.values
```

```
X_tr, X_test, y_tr, y_test = model_selection.train_test_split(X, y, test_size=0.3, ra
# split the train data set into cross validation train and cross validation test
tf_idf_vect = TfidfVectorizer(ngram_range=(1,2))
```

```
X_tr = tf_idf_vect.fit_transform(X_tr)
X_test = tf_idf_vect.transform(X_test)
```

```
print("the type of count vectorizer ",type(X_tr))
print("the shape of out text TFIDF vectorizer ",X_tr.get_shape())
print("the number of unique words including both unigrams and bigrams ", X_tr.get_shap
```

```
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (14000, 305221)
the number of unique words including both unigrams and bigrams 305221
```

```
In [64]: #using truncated svd for dimensionality reduction
```

```
svd = TruncatedSVD(n_components=300)
X_tr=svd.fit_transform(X_tr)
X_test = svd.transform(X_test)
X_tr.shape
```

```
Out[64]: (14000, 300)
```

```
In [65]: #Finding optimal k
```

```
# creating odd list of K for KNN
myList = list(range(5,30))
neighbors = list(filter(lambda x: x % 2 != 0, myList))
```

```
# empty list that will hold cv scores
cv_scores = []
```

```
# perform 2-fold cross validation
```

```
for k in neighbors:
    knn = KNeighborsClassifier(n_neighbors=k,algorithm='kd_tree')
    scores = cross_val_score(knn, X_tr, y_tr, cv=2, scoring='f1')
    cv_scores.append(scores.mean())
```

```
# changing to misclassification error
```

```
MSE = [1 - x for x in cv_scores]
```

```
# determining best k
```

```
optimal_k = neighbors[MSE.index(min(MSE))]
print('\nThe optimal number of neighbors is %d.' % optimal_k)
```

```
# plot misclassification error vs k
```

```
plt.plot(neighbors, MSE)
```

```


for xy in zip(neighbors, np.round(MSE,3)):
    plt.annotate('(%s, %s)' % xy, xy=xy, textcoords='data')

plt.xlabel('Number of Neighbors K')
plt.ylabel('Misclassification Error')
plt.show()

print("the misclassification error for each k value is : ", np.round(MSE,3))

```

The optimal number of neighbors is 23.



the misclassification error for each k value is : [0.095 0.09 0.084 0.082 0.081 0.08 0.079 0.079]

```


In [66]: #Performing KNN using kd_tree
knn = KNeighborsClassifier(n_neighbors=optimal_k,algorithm='kd_tree')
knn.fit(X_tr,y_tr)
pred = knn.predict(X_test)
acc = f1_score(y_test, pred) * float(100)
print('\nThe accuracy of the knn classifier for k = %d is %f%%' % (optimal_k, acc))

```

The accuracy of the knn classifier for k = 23 is 91.877008%

```
In [68]: skplt.metrics.plot_confusion_matrix(y_test ,pred)
```

```
Out[68]: <matplotlib.axes._subplots.AxesSubplot at 0x7f06d7c8bba8>
```



```
In [69]: print(classification_report(y_test ,pred))
```

	precision	recall	f1-score	support
0	0.71	0.12	0.20	950
1	0.86	0.99	0.92	5050
micro avg	0.85	0.85	0.85	6000
macro avg	0.78	0.55	0.56	6000
weighted avg	0.83	0.85	0.80	6000

6 Knn on TFIDF Using Bruteforce

```
In [70]: X = final_data['CleanedText'].values
        y = final_data.Score.values
        X_tr, X_test, y_tr, y_test = model_selection.train_test_split(X, y, test_size=0.3, random_state=42)
        # split the train data set into cross validation train and cross validation test
        tf_idf_vect = TfidfVectorizer(ngram_range=(1,2))

        X_tr = tf_idf_vect.fit_transform(X_tr)
        X_test = tf_idf_vect.transform(X_test)
        print("the type of count vectorizer ",type(X_tr))
        print("the shape of out text TFIDF vectorizer ",X_tr.get_shape())
        print("the number of unique words including both unigrams and bigrams ", X_tr.get_shape()[1])

the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (70000, 993591)
the number of unique words including both unigrams and bigrams 993591
```

```
In [71]: #Finding optimal k

        # creating odd list of K for KNN
        myList = list(range(5,30))
        neighbors = list(filter(lambda x: x % 2 != 0, myList))

        # empty list that will hold cv scores
        cv_scores = []

        # perform 2-fold cross validation
        for k in neighbors:
            knn = KNeighborsClassifier(n_neighbors=k)
            scores = cross_val_score(knn, X_tr, y_tr, cv=2, scoring='f1')
            cv_scores.append(scores.mean())

        # changing to misclassification error
        MSE = [1 - x for x in cv_scores]

        # determining best k
        optimal_k = neighbors[MSE.index(min(MSE))]
        print('\n\nThe optimal number of neighbors is %d.' % optimal_k)

        # plot misclassification error vs k
        plt.plot(neighbors, MSE)


        for xy in zip(neighbors, np.round(MSE,3)):
            plt.annotate('%s, %s' % xy, xy=xy, textcoords='data')

        plt.xlabel('Number of Neighbors K')
```

```
plt.ylabel('Misclassification Error')
plt.show()

print("the misclassification error for each k value is : ", np.round(MSE,3))
```

The optimal number of neighbors is 9.




the misclassification error for each k value is : [0.079 0.078 0.077 0.077 0.077 0.077 0.077 0.077 0.078]

```
In [72]: #Performing KNN using brute force
knn = KNeighborsClassifier(n_neighbors=optimal_k)
knn.fit(X_tr,y_tr)
pred = knn.predict(X_test)
acc = f1_score(y_test, pred) * float(100)
print('\nThe accuracy of the knn classifier for k = %d is %f%%' % (optimal_k, acc))
```

The accuracy of the knn classifier for k = 9 is 92.252721%

```
In [73]: skplt.metrics.plot_confusion_matrix(y_test ,pred)
```

```
Out[73]: <matplotlib.axes._subplots.AxesSubplot at 0x7f06d66821d0>
```



```
In [74]: print(classification_report(y_test ,pred))
```

	precision	recall	f1-score	support
0	0.77	0.13	0.22	4644
1	0.86	0.99	0.92	25356
micro avg	0.86	0.86	0.86	30000
macro avg	0.81	0.56	0.57	30000
weighted avg	0.85	0.86	0.81	30000