

Assignment12

June 1, 2019

```
In [1]: # if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use t
        from keras.utils import np_utils
        from keras.datasets import mnist
        import seaborn as sns
        from keras.initializers import RandomNormal
```

Using TensorFlow backend.

```
In [0]: %matplotlib notebook
        import matplotlib.pyplot as plt
        import numpy as np
        import time
        # https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
        # https://stackoverflow.com/a/14434334
        # this function is used to update the plots for each epoch and error
        def plt_dynamic(x, vy, ty, ax, colors=['b']):
            ax.plot(x, vy, 'b', label="Validation Loss")
            ax.plot(x, ty, 'r', label="Train Loss")
            plt.legend()
            plt.grid()
            fig.canvas.draw()
```

```
In [0]: # the data, shuffled and split between train and test sets
        (X_train, y_train), (X_test, y_test) = mnist.load_data()
```

```
In [4]: print("Number of training examples :", X_train.shape[0], "and each image is of shape (")
        print("Number of testing examples :", X_test.shape[0], "and each image is of shape (%d
```

Number of training examples : 60000 and each image is of shape (28, 28)

Number of testing examples : 10000 and each image is of shape (28, 28)

```
In [0]: # if you observe the input shape its 2 dimensional vector
        # for each image we have a (28*28) vector
        # we will convert the (28*28) vector into single dimensional vector of 1 * 784

        X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
        X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

```
In [6]: X_train.shape
```

```
Out[6]: (60000, 784)
```

```
In [7]: # after converting the input images from 3d to 2d vectors
```

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%)")
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%)")
```

```
Number of training examples : 60000 and each image is of shape (784)
```

```
Number of training examples : 10000 and each image is of shape (784)
```

```
In [8]: # An example data point
```

```
print(X_train[0])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  3  18  18  18 126 136 175  26 166 255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  30  36  94 154
170 253 253 253 253 253 225 172 253 242 195  64  0  0  0  0  0  0
  0  0  0  0  0  49 238 253 253 253 253 253 253 253 251  93  82
 82  56  39  0  0  0  0  0  0  0  0  0  0  0  0  18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  80 156 107 253 253 205  11  0  43 154
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0 14  1 154 253  90  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  139 253 190  2  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  11 190 253  70  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  35 241
225 160 108  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  81 240 253 253 119  25  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  45 186 253 253 150  27  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  16  93 252 253 187
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0 249 253 249  64  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  46 130 183 253
253 207  2  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  39 148 229 253 253 253 250 182  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  24 114 221 253 253 253
253 201  78  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0]
```

0	0	23	66	213	253	253	253	253	198	81	2	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	18	171	219	253	253	253	253	195
80	9	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
55	172	226	253	253	253	253	244	133	11	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	136	253	253	253	212	135	132	16
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0]							

```
In [0]: # if we observe the above matrix each cell is having a value between 0-255
        # before we move to apply machine learning algorithms lets try to normalize the data
        #  $X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$ 
```

```
X_train = X_train/255
X_test = X_test/255
```

```
In [10]: # example data point after normalizing
print(X_train[0])
```

[illegible]

0.	0.	0.01176471	0.07058824	0.07058824	0.07058824
0.49411765	0.53333333	0.68627451	0.10196078	0.65098039	1.
0.96862745	0.49803922	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.11764706	0.14117647	0.36862745	0.60392157
0.66666667	0.99215686	0.99215686	0.99215686	0.99215686	0.99215686
0.88235294	0.6745098	0.99215686	0.94901961	0.76470588	0.25098039
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.19215686
0.93333333	0.99215686	0.99215686	0.99215686	0.99215686	0.99215686
0.99215686	0.99215686	0.99215686	0.98431373	0.36470588	0.32156863
0.32156863	0.21960784	0.15294118	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.07058824	0.85882353	0.99215686
0.99215686	0.99215686	0.99215686	0.99215686	0.77647059	0.71372549
0.96862745	0.94509804	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.31372549	0.61176471	0.41960784	0.99215686
0.99215686	0.80392157	0.04313725	0.	0.16862745	0.60392157
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.05490196	0.00392157	0.60392157	0.99215686	0.35294118
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.54509804	0.99215686	0.74509804	0.00784314	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.04313725
0.74509804	0.99215686	0.2745098	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.1372549	0.94509804
0.88235294	0.62745098	0.42352941	0.00392157	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.31764706	0.94117647	0.99215686
0.99215686	0.46666667	0.09803922	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.17647059	0.72941176	0.99215686	0.99215686

0.58823529	0.10588235	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.0627451	0.36470588	0.98823529	0.99215686	0.73333333
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.97647059	0.99215686	0.97647059	0.25098039	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.18039216	0.50980392	0.71764706	0.99215686
0.99215686	0.81176471	0.00784314	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.15294118	0.58039216
0.89803922	0.99215686	0.99215686	0.99215686	0.98039216	0.71372549
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.09411765	0.44705882	0.86666667	0.99215686	0.99215686	0.99215686
0.99215686	0.78823529	0.30588235	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.09019608	0.25882353	0.83529412	0.99215686
0.99215686	0.99215686	0.99215686	0.77647059	0.31764706	0.00784314
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.07058824	0.67058824
0.85882353	0.99215686	0.99215686	0.99215686	0.99215686	0.76470588
0.31372549	0.03529412	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.21568627	0.6745098	0.88627451	0.99215686	0.99215686	0.99215686
0.99215686	0.95686275	0.52156863	0.04313725	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.53333333	0.99215686
0.99215686	0.99215686	0.83137255	0.52941176	0.51764706	0.0627451
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.

6

```

model.add(Dense(n_layers[0], activation='relu', input_shape=(input_dim,), kernel_initializer=K.he_normal(seed=seed)))
model.add(BatchNormalization())
model.add(Dropout(d[0]))

for i in range(len(n_layers)-1):
    model.add(Dense(n_layers[i+1], activation='relu', kernel_initializer=K.he_normal(seed=seed)))
    model.add(BatchNormalization())
    model.add(Dropout(d[i+1]))

model.add(Dense(output_dim, activation='softmax'))

model.summary()

model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)

%matplotlib inline
score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

#dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

1.1 Architecture with 2 hidden layers.

```
In [78]: n_layers = [468,92]
         d = [0.6,0.5]
         nb_epoch = 25
         nn(n_layers,d)
```

Layer (type)	Output Shape	Param #
dense_128 (Dense)	(None, 468)	367380
batch_normalization_85 (Batch Normalization)	(None, 468)	1872
dropout_65 (Dropout)	(None, 468)	0
dense_129 (Dense)	(None, 92)	43148
batch_normalization_86 (Batch Normalization)	(None, 92)	368
dropout_66 (Dropout)	(None, 92)	0
dense_130 (Dense)	(None, 10)	930

Total params: 413,698

Trainable params: 412,578

Non-trainable params: 1,120

Train on 60000 samples, validate on 10000 samples

Epoch 1/25

60000/60000 [=====] - 13s 210us/step - loss: 0.5165 - acc: 0.8445 - val_loss: 0.2460 - val_acc: 0.9273

Epoch 2/25

60000/60000 [=====] - 4s 64us/step - loss: 0.2460 - acc: 0.9273 - val_loss: 0.1948 - val_acc: 0.9419

Epoch 3/25

60000/60000 [=====] - 4s 65us/step - loss: 0.1948 - acc: 0.9419 - val_loss: 0.1696 - val_acc: 0.9491

Epoch 4/25

60000/60000 [=====] - 4s 65us/step - loss: 0.1696 - acc: 0.9491 - val_loss: 0.1546 - val_acc: 0.9530

Epoch 5/25

60000/60000 [=====] - 4s 65us/step - loss: 0.1546 - acc: 0.9530 - val_loss: 0.1410 - val_acc: 0.9573

Epoch 6/25

60000/60000 [=====] - 4s 66us/step - loss: 0.1410 - acc: 0.9573 - val_loss: 0.1316 - val_acc: 0.9617

Epoch 7/25

60000/60000 [=====] - 4s 65us/step - loss: 0.1316 - acc: 0.9617 - val_loss: 0.1211 - val_acc: 0.9641

Epoch 8/25

60000/60000 [=====] - 4s 65us/step - loss: 0.1211 - acc: 0.9641 - val_loss: 0.1160 - val_acc: 0.9640

Epoch 9/25

60000/60000 [=====] - 4s 66us/step - loss: 0.1160 - acc: 0.9640 - val_loss: 0.1099 - val_acc: 0.9664


Epoch 10/25

60000/60000 [=====] - 4s 65us/step - loss: 0.1099 - acc: 0.9664 - val_loss: 0.1099 - val_acc: 0.9664


```

Epoch 11/25
60000/60000 [=====] - 4s 65us/step - loss: 0.0994 - acc: 0.9702 - val.
Epoch 12/25
60000/60000 [=====] - 4s 65us/step - loss: 0.0975 - acc: 0.9701 - val.
Epoch 13/25
60000/60000 [=====] - 4s 66us/step - loss: 0.0937 - acc: 0.9710 - val.
Epoch 14/25
60000/60000 [=====] - 4s 65us/step - loss: 0.0894 - acc: 0.9720 - val.
Epoch 15/25
60000/60000 [=====] - 4s 64us/step - loss: 0.0891 - acc: 0.9731 - val.
Epoch 16/25
60000/60000 [=====] - 4s 64us/step - loss: 0.0855 - acc: 0.9730 - val.
Epoch 17/25
60000/60000 [=====] - 4s 64us/step - loss: 0.0803 - acc: 0.9752 - val.
Epoch 18/25
60000/60000 [=====] - 4s 64us/step - loss: 0.0812 - acc: 0.9755 - val.
Epoch 19/25
60000/60000 [=====] - 4s 64us/step - loss: 0.0780 - acc: 0.9751 - val.
Epoch 20/25
60000/60000 [=====] - 4s 64us/step - loss: 0.0754 - acc: 0.9772 - val.
Epoch 21/25
60000/60000 [=====] - 4s 64us/step - loss: 0.0730 - acc: 0.9778 - val.
Epoch 22/25
60000/60000 [=====] - 4s 64us/step - loss: 0.0707 - acc: 0.9781 - val.
Epoch 23/25
60000/60000 [=====] - 4s 64us/step - loss: 0.0674 - acc: 0.9791 - val.
Epoch 24/25
60000/60000 [=====] - 4s 64us/step - loss: 0.0701 - acc: 0.9785 - val.
Epoch 25/25
60000/60000 [=====] - 4s 67us/step - loss: 0.0645 - acc: 0.9796 - val.
Test score: 0.058353246974284415
Test accuracy: 0.9828

```



1.2 Architecture with 3 hidden layers.

```
In [79]: n_layers = [512,256,128]
         d = [0.5,0.6,0.5]
         nb_epoch = 25
         nn(n_layers,d)
```

Layer (type)	Output Shape	Param #
dense_131 (Dense)	(None, 512)	401920
batch_normalization_87 (Batch Normalization)	(None, 512)	2048
dropout_67 (Dropout)	(None, 512)	0
dense_132 (Dense)	(None, 256)	131328
batch_normalization_88 (Batch Normalization)	(None, 256)	1024
dropout_68 (Dropout)	(None, 256)	0
dense_133 (Dense)	(None, 128)	32896
batch_normalization_89 (Batch Normalization)	(None, 128)	512

dropout_69 (Dropout) (None, 128) 0

dense_134 (Dense) (None, 10) 1290
=====

Total params: 571,018

Trainable params: 569,226

Non-trainable params: 1,792

Train on 60000 samples, validate on 10000 samples

Epoch 1/25

60000/60000 [=====] - 14s 233us/step - loss: 0.6239 - acc: 0.8107 - val.

Epoch 2/25

60000/60000 [=====] - 5s 78us/step - loss: 0.2689 - acc: 0.9211 - val.

Epoch 3/25

60000/60000 [=====] - 5s 77us/step - loss: 0.2078 - acc: 0.9396 - val.

Epoch 4/25

60000/60000 [=====] - 5s 76us/step - loss: 0.1737 - acc: 0.9488 - val.

Epoch 5/25

60000/60000 [=====] - 5s 77us/step - loss: 0.1565 - acc: 0.9537 - val.

Epoch 6/25

60000/60000 [=====] - 5s 77us/step - loss: 0.1381 - acc: 0.9598 - val.

Epoch 7/25

60000/60000 [=====] - 5s 79us/step - loss: 0.1246 - acc: 0.9625 - val.

Epoch 8/25

60000/60000 [=====] - 5s 78us/step - loss: 0.1201 - acc: 0.9645 - val.

Epoch 9/25

60000/60000 [=====] - 5s 78us/step - loss: 0.1074 - acc: 0.9685 - val.

Epoch 10/25

60000/60000 [=====] - 5s 77us/step - loss: 0.1027 - acc: 0.9692 - val.

Epoch 11/25

60000/60000 [=====] - 5s 79us/step - loss: 0.1010 - acc: 0.9700 - val.

Epoch 12/25

60000/60000 [=====] - 5s 78us/step - loss: 0.0916 - acc: 0.9727 - val.

Epoch 13/25

60000/60000 [=====] - 5s 80us/step - loss: 0.0882 - acc: 0.9732 - val.

Epoch 14/25

60000/60000 [=====] - 5s 79us/step - loss: 0.0857 - acc: 0.9744 - val.

Epoch 15/25

60000/60000 [=====] - 5s 82us/step - loss: 0.0832 - acc: 0.9750 - val.

Epoch 16/25

60000/60000 [=====] - 5s 80us/step - loss: 0.0758 - acc: 0.9770 - val.

Epoch 17/25

60000/60000 [=====] - 5s 80us/step - loss: 0.0717 - acc: 0.9786 - val.

Epoch 18/25

60000/60000 [=====] - 5s 82us/step - loss: 0.0724 - acc: 0.9783 - val.


Epoch 19/25

60000/60000 [=====] - 5s 81us/step - loss: 0.0684 - acc: 0.9798 - val.

```

Epoch 20/25
60000/60000 [=====] - 5s 81us/step - loss: 0.0708 - acc: 0.9790 - val.
Epoch 21/25
60000/60000 [=====] - 5s 79us/step - loss: 0.0657 - acc: 0.9802 - val.
Epoch 22/25
60000/60000 [=====] - 5s 80us/step - loss: 0.0627 - acc: 0.9810 - val.
Epoch 23/25
60000/60000 [=====] - 5s 79us/step - loss: 0.0620 - acc: 0.9813 - val.
Epoch 24/25
60000/60000 [=====] - 5s 79us/step - loss: 0.0614 - acc: 0.9813 - val.
Epoch 25/25
60000/60000 [=====] - 5s 80us/step - loss: 0.0612 - acc: 0.9815 - val.
Test score: 0.0548867928153486
Test accuracy: 0.985

```



1.3 Architecture with 4 hidden layers(for experimentation)

```

In [56]: #Using 4 hidden layers for experimentation
n_layers = [600,460,290,112]
d = [0.4,0.7,0.5,0.7]
nb_epoch = 25
nn(n_layers,d)

```

Layer (type)	Output Shape	Param #
--------------	--------------	---------

dense_56 (Dense)	(None, 600)	471000
batch_normalization_44 (Batch Normalization)	(None, 600)	2400
dropout_44 (Dropout)	(None, 600)	0
dense_57 (Dense)	(None, 460)	276460
batch_normalization_45 (Batch Normalization)	(None, 460)	1840
dropout_45 (Dropout)	(None, 460)	0
dense_58 (Dense)	(None, 290)	133690
batch_normalization_46 (Batch Normalization)	(None, 290)	1160
dropout_46 (Dropout)	(None, 290)	0
dense_59 (Dense)	(None, 112)	32592
batch_normalization_47 (Batch Normalization)	(None, 112)	448
dropout_47 (Dropout)	(None, 112)	0
dense_60 (Dense)	(None, 10)	1130

Total params: 920,720
 Trainable params: 917,796
 Non-trainable params: 2,924

Train on 60000 samples, validate on 10000 samples

Epoch 1/25

60000/60000 [=====] - 10s 175us/step - loss: 1.0535 - acc: 0.6908 - val_loss: 1.0535 - val_acc: 0.6908

Epoch 2/25

60000/60000 [=====] - 5s 91us/step - loss: 0.3588 - acc: 0.8999 - val_loss: 0.3588 - val_acc: 0.8999

Epoch 3/25

60000/60000 [=====] - 5s 89us/step - loss: 0.2583 - acc: 0.9291 - val_loss: 0.2583 - val_acc: 0.9291

Epoch 4/25

60000/60000 [=====] - 5s 90us/step - loss: 0.2100 - acc: 0.9431 - val_loss: 0.2100 - val_acc: 0.9431

Epoch 5/25

60000/60000 [=====] - 5s 90us/step - loss: 0.1803 - acc: 0.9518 - val_loss: 0.1803 - val_acc: 0.9518

Epoch 6/25

60000/60000 [=====] - 5s 91us/step - loss: 0.1652 - acc: 0.9559 - val_loss: 0.1652 - val_acc: 0.9559

Epoch 7/25

60000/60000 [=====] - 5s 89us/step - loss: 0.1473 - acc: 0.9612 - val_loss: 0.1473 - val_acc: 0.9612


Epoch 8/25

60000/60000 [=====] - 5s 90us/step - loss: 0.1353 - acc: 0.9637 - val_loss: 0.1353 - val_acc: 0.9637

```

Epoch 9/25
60000/60000 [=====] - 5s 89us/step - loss: 0.1268 - acc: 0.9655 - val.
Epoch 10/25
60000/60000 [=====] - 5s 90us/step - loss: 0.1123 - acc: 0.9691 - val.
Epoch 11/25
60000/60000 [=====] - 5s 90us/step - loss: 0.1104 - acc: 0.9708 - val.
Epoch 12/25
60000/60000 [=====] - 5s 90us/step - loss: 0.1042 - acc: 0.9724 - val.
Epoch 13/25
60000/60000 [=====] - 5s 90us/step - loss: 0.1020 - acc: 0.9728 - val.
Epoch 14/25
60000/60000 [=====] - 5s 90us/step - loss: 0.0949 - acc: 0.9748 - val.
Epoch 15/25
60000/60000 [=====] - 5s 91us/step - loss: 0.0911 - acc: 0.9756 - val.
Epoch 16/25
60000/60000 [=====] - 6s 92us/step - loss: 0.0876 - acc: 0.9763 - val.
Epoch 17/25
60000/60000 [=====] - 5s 90us/step - loss: 0.0857 - acc: 0.9771 - val.
Epoch 18/25
60000/60000 [=====] - 5s 89us/step - loss: 0.0789 - acc: 0.9784 - val.
Epoch 19/25
60000/60000 [=====] - 5s 89us/step - loss: 0.0773 - acc: 0.9789 - val.
Epoch 20/25
60000/60000 [=====] - 5s 89us/step - loss: 0.0749 - acc: 0.9799 - val.
Epoch 21/25
60000/60000 [=====] - 5s 89us/step - loss: 0.0741 - acc: 0.9804 - val.
Epoch 22/25
60000/60000 [=====] - 5s 89us/step - loss: 0.0678 - acc: 0.9815 - val.
Epoch 23/25
60000/60000 [=====] - 5s 89us/step - loss: 0.0638 - acc: 0.9824 - val.
Epoch 24/25
60000/60000 [=====] - 5s 91us/step - loss: 0.0631 - acc: 0.9827 - val.
Epoch 25/25
60000/60000 [=====] - 6s 93us/step - loss: 0.0612 - acc: 0.9831 - val.
Test score: 0.06384382131070597
Test accuracy: 0.9847

```



1.4 Architecture with 5 hidden layers.

```
In [59]: n_layers = [650,520,410,290,125]
         d = [0.5,0.6,0.5,0.5,0.5]
         nb_epoch = 25
         nn(n_layers,d)
```

Layer (type)	Output Shape	Param #
dense_72 (Dense)	(None, 650)	510250
batch_normalization_58 (Batch Normalization)	(None, 650)	2600
dropout_57 (Dropout)	(None, 650)	0
dense_73 (Dense)	(None, 520)	338520
batch_normalization_59 (Batch Normalization)	(None, 520)	2080
dropout_58 (Dropout)	(None, 520)	0
dense_74 (Dense)	(None, 410)	213610
batch_normalization_60 (Batch Normalization)	(None, 410)	1640

dropout_59 (Dropout)	(None, 410)	0
dense_75 (Dense)	(None, 290)	119190
batch_normalization_61 (Batch Normalization)	(None, 290)	1160
dropout_60 (Dropout)	(None, 290)	0
dense_76 (Dense)	(None, 125)	36375
batch_normalization_62 (Batch Normalization)	(None, 125)	500
dropout_61 (Dropout)	(None, 125)	0
dense_77 (Dense)	(None, 10)	1260

=====
Total params: 1,227,185

Trainable params: 1,223,195

Non-trainable params: 3,990

Train on 60000 samples, validate on 10000 samples

Epoch 1/25

60000/60000 [=====] - 13s 215us/step - loss: 0.9056 - acc: 0.7211 - val_loss: 0.9056 - val_acc: 0.7211

Epoch 2/25

60000/60000 [=====] - 6s 107us/step - loss: 0.3179 - acc: 0.9073 - val_loss: 0.3179 - val_acc: 0.9073

Epoch 3/25

60000/60000 [=====] - 6s 105us/step - loss: 0.2403 - acc: 0.9319 - val_loss: 0.2403 - val_acc: 0.9319

Epoch 4/25

60000/60000 [=====] - 6s 107us/step - loss: 0.1951 - acc: 0.9447 - val_loss: 0.1951 - val_acc: 0.9447

Epoch 5/25

60000/60000 [=====] - 7s 110us/step - loss: 0.1753 - acc: 0.9509 - val_loss: 0.1753 - val_acc: 0.9509

Epoch 6/25

60000/60000 [=====] - 6s 107us/step - loss: 0.1575 - acc: 0.9554 - val_loss: 0.1575 - val_acc: 0.9554

Epoch 7/25

60000/60000 [=====] - 6s 106us/step - loss: 0.1490 - acc: 0.9583 - val_loss: 0.1490 - val_acc: 0.9583

Epoch 8/25

60000/60000 [=====] - 6s 105us/step - loss: 0.1349 - acc: 0.9621 - val_loss: 0.1349 - val_acc: 0.9621

Epoch 9/25

60000/60000 [=====] - 6s 106us/step - loss: 0.1284 - acc: 0.9640 - val_loss: 0.1284 - val_acc: 0.9640

Epoch 10/25

60000/60000 [=====] - 6s 105us/step - loss: 0.1186 - acc: 0.9668 - val_loss: 0.1186 - val_acc: 0.9668

Epoch 11/25

60000/60000 [=====] - 6s 107us/step - loss: 0.1163 - acc: 0.9672 - val_loss: 0.1163 - val_acc: 0.9672


Epoch 12/25

60000/60000 [=====] - 6s 106us/step - loss: 0.1090 - acc: 0.9697 - val_loss: 0.1090 - val_acc: 0.9697

Epoch 13/25

60000/60000 [=====] - 6s 105us/step - loss: 0.1050 - acc: 0.9705 - val_loss: 0.1050 - val_acc: 0.9705

Epoch 14/25
60000/60000 [=====] - 6s 106us/step - loss: 0.1008 - acc: 0.9720 - va.
Epoch 15/25
60000/60000 [=====] - 6s 107us/step - loss: 0.0978 - acc: 0.9729 - va.
Epoch 16/25
60000/60000 [=====] - 7s 116us/step - loss: 0.0923 - acc: 0.9743 - va.
Epoch 17/25
60000/60000 [=====] - 7s 111us/step - loss: 0.0884 - acc: 0.9753 - va.
Epoch 18/25
60000/60000 [=====] - 7s 110us/step - loss: 0.0896 - acc: 0.9744 - va.
Epoch 19/25
60000/60000 [=====] - 6s 108us/step - loss: 0.0825 - acc: 0.9768 - va.
Epoch 20/25
60000/60000 [=====] - 6s 108us/step - loss: 0.0803 - acc: 0.9782 - va.
Epoch 21/25
60000/60000 [=====] - 6s 107us/step - loss: 0.0735 - acc: 0.9790 - va.
Epoch 22/25
60000/60000 [=====] - 6s 107us/step - loss: 0.0722 - acc: 0.9789 - va.
Epoch 23/25
60000/60000 [=====] - 6s 108us/step - loss: 0.0738 - acc: 0.9789 - va.
Epoch 24/25
60000/60000 [=====] - 6s 108us/step - loss: 0.0706 - acc: 0.9795 - va.
Epoch 25/25
60000/60000 [=====] - 6s 107us/step - loss: 0.0669 - acc: 0.9807 - va.
Test score: 0.06073720659725368
Test accuracy: 0.9846



2 Defining a function for model with only batch normalisation and no dropout.

In [0]: # <https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization->

```
def nn1(n_layers,bn):
    model = Sequential()
    model.add(Dense(n_layers[0], activation='relu', input_shape=(input_dim,), kernel_initializer=K.he_normal()))
    if(bn==1):
        model.add(BatchNormalization())
        for i in range(len(n_layers)-1):
            model.add(Dense(n_layers[i+1], activation='relu',kernel_initializer=K.he_normal()))
            model.add(BatchNormalization())
    else:
        for i in range(len(n_layers)-1):
            model.add(Dense(n_layers[i+1], activation='relu',kernel_initializer=K.he_normal()))

    model.add(Dense(output_dim, activation='softmax'))

    model.summary()

    model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

    history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)

    %matplotlib inline
    score = model.evaluate(X_test, Y_test, verbose=0)
    print('Test score:', score[0])
    print('Test accuracy:', score[1])

    fig,ax = plt.subplots(1,1)
    ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

    # list of epoch numbers
    x = list(range(1,nb_epoch+1))

    #dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
    # history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,
```

```

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

2.0.1 Architecture with 2 hidden layers.

```

In [70]: n_layers = [468,92]
         bn = 1
         nb_epoch = 25
         nn1(n_layers,bn)

```

Layer (type)	Output Shape	Param #
dense_102 (Dense)	(None, 468)	367380
batch_normalization_75 (Batch Normalization)	(None, 468)	1872
dense_103 (Dense)	(None, 92)	43148
batch_normalization_76 (Batch Normalization)	(None, 92)	368
dense_104 (Dense)	(None, 10)	930

```

Total params: 413,698
Trainable params: 412,578
Non-trainable params: 1,120

```

```

Train on 60000 samples, validate on 10000 samples

```

```


Epoch 1/25
60000/60000 [=====] - 11s 179us/step - loss: 0.1946 - acc: 0.9427 - val_loss: 0.0741 - val_acc: 0.9776
Epoch 2/25
60000/60000 [=====] - 4s 62us/step - loss: 0.0741 - acc: 0.9776 - val_loss: 0.0479 - val_acc: 0.9853
Epoch 3/25
60000/60000 [=====] - 4s 63us/step - loss: 0.0479 - acc: 0.9853 - val_loss: 0.0355 - val_acc: 0.9889
Epoch 4/25
60000/60000 [=====] - 4s 63us/step - loss: 0.0355 - acc: 0.9889 - val_loss: 0.0275 - val_acc: 0.9914
Epoch 5/25
60000/60000 [=====] - 4s 63us/step - loss: 0.0275 - acc: 0.9914 - val_loss: 0.0224 - val_acc: 0.9926
Epoch 6/25
60000/60000 [=====] - 4s 64us/step - loss: 0.0224 - acc: 0.9926 - val_loss: 0.0224 - val_acc: 0.9926

```

```

Epoch 7/25
60000/60000 [=====] - 4s 64us/step - loss: 0.0187 - acc: 0.9942 - val.
Epoch 8/25
60000/60000 [=====] - 4s 64us/step - loss: 0.0158 - acc: 0.9947 - val.
Epoch 9/25
60000/60000 [=====] - 4s 66us/step - loss: 0.0157 - acc: 0.9948 - val.
Epoch 10/25
60000/60000 [=====] - 4s 63us/step - loss: 0.0129 - acc: 0.9959 - val.
Epoch 11/25
60000/60000 [=====] - 4s 62us/step - loss: 0.0142 - acc: 0.9951 - val.
Epoch 12/25
60000/60000 [=====] - 4s 64us/step - loss: 0.0107 - acc: 0.9965 - val.
Epoch 13/25
60000/60000 [=====] - 4s 63us/step - loss: 0.0104 - acc: 0.9966 - val.
Epoch 14/25
60000/60000 [=====] - 4s 62us/step - loss: 0.0108 - acc: 0.9963 - val.
Epoch 15/25
60000/60000 [=====] - 4s 63us/step - loss: 0.0107 - acc: 0.9963 - val.
Epoch 16/25
60000/60000 [=====] - 4s 61us/step - loss: 0.0085 - acc: 0.9971 - val.
Epoch 17/25
60000/60000 [=====] - 4s 63us/step - loss: 0.0071 - acc: 0.9977 - val.
Epoch 18/25
60000/60000 [=====] - 4s 62us/step - loss: 0.0068 - acc: 0.9978 - val.
Epoch 19/25
60000/60000 [=====] - 4s 62us/step - loss: 0.0103 - acc: 0.9966 - val.
Epoch 20/25
60000/60000 [=====] - 4s 62us/step - loss: 0.0085 - acc: 0.9971 - val.
Epoch 21/25
60000/60000 [=====] - 4s 62us/step - loss: 0.0064 - acc: 0.9978 - val.
Epoch 22/25
60000/60000 [=====] - 4s 62us/step - loss: 0.0063 - acc: 0.9981 - val.
Epoch 23/25
60000/60000 [=====] - 4s 63us/step - loss: 0.0069 - acc: 0.9978 - val.
Epoch 24/25
60000/60000 [=====] - 4s 62us/step - loss: 0.0050 - acc: 0.9983 - val.
Epoch 25/25
60000/60000 [=====] - 4s 62us/step - loss: 0.0050 - acc: 0.9984 - val.
Test score: 0.08396656601358045
Test accuracy: 0.9808

```



2.1 Architecture with 3 hidden layers.

```
In [71]: n_layers = [512,256,128]
         bn=1
         nb_epoch = 25
         nn1(n_layers,bn)
```


Layer (type)	Output Shape	Param #
dense_105 (Dense)	(None, 512)	401920
batch_normalization_77 (Batch Normalization)	(None, 512)	2048
dense_106 (Dense)	(None, 256)	131328
batch_normalization_78 (Batch Normalization)	(None, 256)	1024
dense_107 (Dense)	(None, 128)	32896
batch_normalization_79 (Batch Normalization)	(None, 128)	512
dense_108 (Dense)	(None, 10)	1290
Total params: 571,018		

Trainable params: 569,226
Non-trainable params: 1,792

Train on 60000 samples, validate on 10000 samples

Epoch 1/25
60000/60000 [=====] - 12s 206us/step - loss: 0.1840 - acc: 0.9448 - val.
Epoch 2/25
60000/60000 [=====] - 5s 75us/step - loss: 0.0713 - acc: 0.9782 - val.
Epoch 3/25
60000/60000 [=====] - 5s 76us/step - loss: 0.0490 - acc: 0.9843 - val.
Epoch 4/25
60000/60000 [=====] - 5s 78us/step - loss: 0.0337 - acc: 0.9889 - val.
Epoch 5/25
60000/60000 [=====] - 5s 78us/step - loss: 0.0315 - acc: 0.9896 - val.
Epoch 6/25
60000/60000 [=====] - 5s 75us/step - loss: 0.0259 - acc: 0.9912 - val.
Epoch 7/25
60000/60000 [=====] - 4s 74us/step - loss: 0.0225 - acc: 0.9926 - val.
Epoch 8/25
60000/60000 [=====] - 4s 74us/step - loss: 0.0221 - acc: 0.9928 - val.
Epoch 9/25
60000/60000 [=====] - 4s 74us/step - loss: 0.0177 - acc: 0.9940 - val.
Epoch 10/25
60000/60000 [=====] - 4s 74us/step - loss: 0.0149 - acc: 0.9950 - val.
Epoch 11/25
60000/60000 [=====] - 5s 75us/step - loss: 0.0144 - acc: 0.9955 - val.
Epoch 12/25
60000/60000 [=====] - 4s 74us/step - loss: 0.0121 - acc: 0.9963 - val.
Epoch 13/25
60000/60000 [=====] - 4s 75us/step - loss: 0.0140 - acc: 0.9953 - val.
Epoch 14/25
60000/60000 [=====] - 4s 74us/step - loss: 0.0136 - acc: 0.9956 - val.
Epoch 15/25
60000/60000 [=====] - 5s 76us/step - loss: 0.0128 - acc: 0.9957 - val.
Epoch 16/25
60000/60000 [=====] - 5s 79us/step - loss: 0.0090 - acc: 0.9970 - val.
Epoch 17/25
60000/60000 [=====] - 4s 74us/step - loss: 0.0100 - acc: 0.9965 - val.
Epoch 18/25
60000/60000 [=====] - 5s 76us/step - loss: 0.0103 - acc: 0.9969 - val.
Epoch 19/25
60000/60000 [=====] - 5s 76us/step - loss: 0.0086 - acc: 0.9973 - val.
Epoch 20/25
60000/60000 [=====] - 5s 77us/step - loss: 0.0072 - acc: 0.9977 - val.
Epoch 21/25
60000/60000 [=====] - 4s 75us/step - loss: 0.0101 - acc: 0.9962 - val.
Epoch 22/25
60000/60000 [=====] - 4s 74us/step - loss: 0.0103 - acc: 0.9967 - val.

Epoch 23/25
60000/60000 [=====] - 4s 75us/step - loss: 0.0069 - acc: 0.9977 - val.
Epoch 24/25
60000/60000 [=====] - 4s 75us/step - loss: 0.0060 - acc: 0.9981 - val.
Epoch 25/25
60000/60000 [=====] - 4s 75us/step - loss: 0.0058 - acc: 0.9981 - val.
Test score: 0.07232702655499207
Test accuracy: 0.9833



2.2 Architecture with 5 hidden layers.

```
In [72]: n_layers = [650,520,410,290,125]
         bn=1
         nb_epoch = 25
         nn1(n_layers,bn)
```

Layer (type)	Output Shape	Param #
dense_109 (Dense)	(None, 650)	510250
batch_normalization_80 (Batc	(None, 650)	2600
dense_110 (Dense)	(None, 520)	338520

batch_normalization_81 (Batch Normalization)	(None, 520)	2080
dense_111 (Dense)	(None, 410)	213610
batch_normalization_82 (Batch Normalization)	(None, 410)	1640
dense_112 (Dense)	(None, 290)	119190
batch_normalization_83 (Batch Normalization)	(None, 290)	1160
dense_113 (Dense)	(None, 125)	36375
batch_normalization_84 (Batch Normalization)	(None, 125)	500
dense_114 (Dense)	(None, 10)	1260

=====
Total params: 1,227,185

Trainable params: 1,223,195

Non-trainable params: 3,990

Train on 60000 samples, validate on 10000 samples

Epoch 1/25

60000/60000 [=====] - 15s 242us/step - loss: 0.1877 - acc: 0.9439 - val_loss: 0.1877 - val_acc: 0.9439

Epoch 2/25

60000/60000 [=====] - 6s 102us/step - loss: 0.0808 - acc: 0.9746 - val_loss: 0.0808 - val_acc: 0.9746

Epoch 3/25

60000/60000 [=====] - 6s 100us/step - loss: 0.0581 - acc: 0.9813 - val_loss: 0.0581 - val_acc: 0.9813

Epoch 4/25

60000/60000 [=====] - 6s 103us/step - loss: 0.0460 - acc: 0.9852 - val_loss: 0.0460 - val_acc: 0.9852

Epoch 5/25

60000/60000 [=====] - 6s 103us/step - loss: 0.0400 - acc: 0.9869 - val_loss: 0.0400 - val_acc: 0.9869

Epoch 6/25

60000/60000 [=====] - 6s 104us/step - loss: 0.0360 - acc: 0.9882 - val_loss: 0.0360 - val_acc: 0.9882

Epoch 7/25

60000/60000 [=====] - 6s 104us/step - loss: 0.0323 - acc: 0.9894 - val_loss: 0.0323 - val_acc: 0.9894

Epoch 8/25

60000/60000 [=====] - 6s 104us/step - loss: 0.0270 - acc: 0.9909 - val_loss: 0.0270 - val_acc: 0.9909

Epoch 9/25

60000/60000 [=====] - 6s 102us/step - loss: 0.0256 - acc: 0.9915 - val_loss: 0.0256 - val_acc: 0.9915

Epoch 10/25

60000/60000 [=====] - 6s 102us/step - loss: 0.0220 - acc: 0.9927 - val_loss: 0.0220 - val_acc: 0.9927

Epoch 11/25

60000/60000 [=====] - 6s 103us/step - loss: 0.0245 - acc: 0.9918 - val_loss: 0.0245 - val_acc: 0.9918


Epoch 12/25

60000/60000 [=====] - 6s 102us/step - loss: 0.0215 - acc: 0.9928 - val_loss: 0.0215 - val_acc: 0.9928

Epoch 13/25

60000/60000 [=====] - 6s 102us/step - loss: 0.0221 - acc: 0.9928 - val_loss: 0.0221 - val_acc: 0.9928

Epoch 14/25
 60000/60000 [=====] - 6s 101us/step - loss: 0.0140 - acc: 0.9954 - va.
 Epoch 15/25
 60000/60000 [=====] - 6s 101us/step - loss: 0.0177 - acc: 0.9939 - va.
 Epoch 16/25
 60000/60000 [=====] - 6s 101us/step - loss: 0.0178 - acc: 0.9941 - va.
 Epoch 17/25
 60000/60000 [=====] - 6s 102us/step - loss: 0.0124 - acc: 0.9960 - va.
 Epoch 18/25
 60000/60000 [=====] - 6s 102us/step - loss: 0.0142 - acc: 0.9955 - va.
 Epoch 19/25
 60000/60000 [=====] - 6s 102us/step - loss: 0.0140 - acc: 0.9953 - va.
 Epoch 20/25
 60000/60000 [=====] - 6s 103us/step - loss: 0.0139 - acc: 0.9953 - va.
 Epoch 21/25
 60000/60000 [=====] - 6s 104us/step - loss: 0.0121 - acc: 0.9957 - va.
 Epoch 22/25
 60000/60000 [=====] - 6s 100us/step - loss: 0.0104 - acc: 0.9966 - va.
 Epoch 23/25
 60000/60000 [=====] - 6s 101us/step - loss: 0.0113 - acc: 0.9963 - va.
 Epoch 24/25
 60000/60000 [=====] - 6s 99us/step - loss: 0.0116 - acc: 0.9962 - va.
 Epoch 25/25
 60000/60000 [=====] - 6s 101us/step - loss: 0.0093 - acc: 0.9972 - va.
 Test score: 0.06976913923782413
 Test accuracy: 0.984



3 Defining a function for model with no dropout and no batch normalization.

In [0]: # <https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization->

```
def nn2(n_layers):
    model = Sequential()
    model.add(Dense(n_layers[0], activation='relu', input_shape=(input_dim,), kernel_initializer=K.he_normal(
    for i in range(len(n_layers)-1):
        model.add(Dense(n_layers[i+1], activation='relu', kernel_initializer=K.he_normal(

    model.add(Dense(output_dim, activation='softmax'))

    model.summary()

    model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

    history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)

    %matplotlib inline
    score = model.evaluate(X_test, Y_test, verbose=0)
    print('Test score:', score[0])
    print('Test accuracy:', score[1])

    fig, ax = plt.subplots(1, 1)
    ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

    # list of epoch numbers
    x = list(range(1, nb_epoch+1))

    # dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
    # history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,

    # we will get val_loss and val_acc only when you pass the paramter validation_data
    # val_loss : validation loss
    # val_acc : validation accuracy

    # loss : training loss
    # acc : train accuracy
```

```

# for each key in history.history we will have a list of length equal to number of
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

3.0.1 Architecture with 2 hidden layers.

```

In [75]: n_layers = [468,92]
         nb_epoch = 25
         nn2(n_layers)

```

Layer (type)	Output Shape	Param #
dense_115 (Dense)	(None, 468)	367380
dense_116 (Dense)	(None, 92)	43148
dense_117 (Dense)	(None, 10)	930

Total params: 411,458
 Trainable params: 411,458
 Non-trainable params: 0


Train on 60000 samples, validate on 10000 samples

Epoch	60000/60000	loss	acc	val_loss	val_acc
Epoch 1/25	60000/60000	0.2425	0.9295	0.2425	0.9295
Epoch 2/25	60000/60000	0.0907	0.9724	0.0907	0.9724
Epoch 3/25	60000/60000	0.0576	0.9822	0.0576	0.9822
Epoch 4/25	60000/60000	0.0402	0.9878	0.0402	0.9878
Epoch 5/25	60000/60000	0.0292	0.9905	0.0292	0.9905
Epoch 6/25	60000/60000	0.0203	0.9936	0.0203	0.9936
Epoch 7/25	60000/60000	0.0169	0.9944	0.0169	0.9944
Epoch 8/25	60000/60000	0.0148	0.9951	0.0148	0.9951
Epoch 9/25	60000/60000	0.0126	0.9960	0.0126	0.9960
Epoch 10/25	60000/60000	0.0124	0.9958	0.0124	0.9958
Epoch 11/25	60000/60000	0.0095	0.9969	0.0095	0.9969
Epoch 12/25	60000/60000				

```

60000/60000 [=====] - 3s 44us/step - loss: 0.0097 - acc: 0.9968 - val.
Epoch 13/25
60000/60000 [=====] - 3s 44us/step - loss: 0.0089 - acc: 0.9971 - val.
Epoch 14/25
60000/60000 [=====] - 3s 44us/step - loss: 0.0093 - acc: 0.9969 - val.
Epoch 15/25
60000/60000 [=====] - 3s 42us/step - loss: 0.0089 - acc: 0.9971 - val.
Epoch 16/25
60000/60000 [=====] - 3s 43us/step - loss: 0.0064 - acc: 0.9980 - val.
Epoch 17/25
60000/60000 [=====] - 3s 42us/step - loss: 0.0044 - acc: 0.9986 - val.
Epoch 18/25
60000/60000 [=====] - 2s 41us/step - loss: 0.0056 - acc: 0.9981 - val.
Epoch 19/25
60000/60000 [=====] - 3s 42us/step - loss: 0.0101 - acc: 0.9967 - val.
Epoch 20/25
60000/60000 [=====] - 3s 42us/step - loss: 0.0055 - acc: 0.9979 - val.
Epoch 21/25
60000/60000 [=====] - 3s 42us/step - loss: 0.0055 - acc: 0.9981 - val.
Epoch 22/25
60000/60000 [=====] - 3s 42us/step - loss: 0.0074 - acc: 0.9978 - val.
Epoch 23/25
60000/60000 [=====] - 2s 42us/step - loss: 0.0042 - acc: 0.9986 - val.
Epoch 24/25
60000/60000 [=====] - 3s 42us/step - loss: 0.0082 - acc: 0.9974 - val.
Epoch 25/25
60000/60000 [=====] - 3s 42us/step - loss: 0.0065 - acc: 0.9978 - val.
Test score: 0.09983180044603286
Test accuracy: 0.9807

```



3.1 Architecture with 3 hidden layers

```
In [76]: n_layers = [512,256,128]
         nb_epoch = 25
         nn2(n_layers)
```

Layer (type)	Output Shape	Param #
dense_118 (Dense)	(None, 512)	401920
dense_119 (Dense)	(None, 256)	131328
dense_120 (Dense)	(None, 128)	32896
dense_121 (Dense)	(None, 10)	1290

```
Total params: 567,434
Trainable params: 567,434
Non-trainable params: 0
```

```
-----
Train on 60000 samples, validate on 10000 samples
```

```
Epoch 1/25
```

```
60000/60000 [=====] - 11s 180us/step - loss: 0.2199 - acc: 0.9339 - v
```


```
Epoch 2/25
```

```

60000/60000 [=====] - 3s 46us/step - loss: 0.0813 - acc: 0.9754 - val.
Epoch 3/25
60000/60000 [=====] - 3s 45us/step - loss: 0.0528 - acc: 0.9829 - val.
Epoch 4/25
60000/60000 [=====] - 3s 44us/step - loss: 0.0376 - acc: 0.9877 - val.
Epoch 5/25
60000/60000 [=====] - 3s 45us/step - loss: 0.0279 - acc: 0.9910 - val.
Epoch 6/25
60000/60000 [=====] - 3s 45us/step - loss: 0.0254 - acc: 0.9915 - val.
Epoch 7/25
60000/60000 [=====] - 3s 45us/step - loss: 0.0205 - acc: 0.9932 - val.
Epoch 8/25
60000/60000 [=====] - 3s 45us/step - loss: 0.0166 - acc: 0.9945 - val.
Epoch 9/25
60000/60000 [=====] - 3s 45us/step - loss: 0.0181 - acc: 0.9941 - val.
Epoch 10/25
60000/60000 [=====] - 3s 45us/step - loss: 0.0171 - acc: 0.9942 - val.
Epoch 11/25
60000/60000 [=====] - 3s 46us/step - loss: 0.0135 - acc: 0.9956 - val.
Epoch 12/25
60000/60000 [=====] - 3s 46us/step - loss: 0.0135 - acc: 0.9957 - val.
Epoch 13/25
60000/60000 [=====] - 3s 47us/step - loss: 0.0116 - acc: 0.9965 - val.
Epoch 14/25
60000/60000 [=====] - 3s 47us/step - loss: 0.0130 - acc: 0.9957 - val.
Epoch 15/25
60000/60000 [=====] - 3s 46us/step - loss: 0.0095 - acc: 0.9970 - val.
Epoch 16/25
60000/60000 [=====] - 3s 45us/step - loss: 0.0111 - acc: 0.9965 - val.
Epoch 17/25
60000/60000 [=====] - 3s 45us/step - loss: 0.0097 - acc: 0.9970 - val.
Epoch 18/25
60000/60000 [=====] - 3s 45us/step - loss: 0.0080 - acc: 0.9974 - val.
Epoch 19/25
60000/60000 [=====] - 3s 45us/step - loss: 0.0106 - acc: 0.9968 - val.
Epoch 20/25
60000/60000 [=====] - 3s 45us/step - loss: 0.0072 - acc: 0.9980 - val.
Epoch 21/25
60000/60000 [=====] - 3s 45us/step - loss: 0.0082 - acc: 0.9976 - val.
Epoch 22/25
60000/60000 [=====] - 3s 45us/step - loss: 0.0088 - acc: 0.9974 - val.
Epoch 23/25
60000/60000 [=====] - 3s 45us/step - loss: 0.0065 - acc: 0.9981 - val.
Epoch 24/25
60000/60000 [=====] - 3s 45us/step - loss: 0.0093 - acc: 0.9973 - val.
Epoch 25/25
60000/60000 [=====] - 3s 45us/step - loss: 0.0072 - acc: 0.9979 - val.
Test score: 0.0932893890325599

```

Test accuracy: 0.9814



3.2 Architecture with 5 hidden layers.

```
In [77]: n_layers = [650,520,410,290,125]
         nb_epoch = 25
         nn2(n_layers)
```

Layer (type)	Output Shape	Param #
dense_122 (Dense)	(None, 650)	510250
dense_123 (Dense)	(None, 520)	338520
dense_124 (Dense)	(None, 410)	213610
dense_125 (Dense)	(None, 290)	119190
dense_126 (Dense)	(None, 125)	36375
dense_127 (Dense)	(None, 10)	1260
Total params: 1,219,205		

Trainable params: 1,219,205

Non-trainable params: 0

Train on 60000 samples, validate on 10000 samples

Epoch 1/25

60000/60000 [=====] - 11s 190us/step - loss: 0.2114 - acc: 0.9360 - val.

Epoch 2/25

60000/60000 [=====] - 3s 53us/step - loss: 0.0887 - acc: 0.9726 - val.

Epoch 3/25

60000/60000 [=====] - 3s 53us/step - loss: 0.0605 - acc: 0.9813 - val.

Epoch 4/25

60000/60000 [=====] - 3s 52us/step - loss: 0.0463 - acc: 0.9858 - val.

Epoch 5/25

60000/60000 [=====] - 3s 52us/step - loss: 0.0388 - acc: 0.9880 - val.

Epoch 6/25

60000/60000 [=====] - 3s 53us/step - loss: 0.0342 - acc: 0.9895 - val.

Epoch 7/25

60000/60000 [=====] - 3s 52us/step - loss: 0.0268 - acc: 0.9914 - val.

Epoch 8/25

60000/60000 [=====] - 3s 53us/step - loss: 0.0274 - acc: 0.9916 - val.

Epoch 9/25

60000/60000 [=====] - 3s 52us/step - loss: 0.0248 - acc: 0.9923 - val.

Epoch 10/25

60000/60000 [=====] - 3s 52us/step - loss: 0.0196 - acc: 0.9942 - val.

Epoch 11/25

60000/60000 [=====] - 3s 52us/step - loss: 0.0189 - acc: 0.9944 - val.

Epoch 12/25

60000/60000 [=====] - 3s 54us/step - loss: 0.0167 - acc: 0.9952 - val.

Epoch 13/25

60000/60000 [=====] - 3s 53us/step - loss: 0.0177 - acc: 0.9949 - val.

Epoch 14/25

60000/60000 [=====] - 3s 53us/step - loss: 0.0155 - acc: 0.9952 - val.

Epoch 15/25

60000/60000 [=====] - 3s 54us/step - loss: 0.0170 - acc: 0.9951 - val.

Epoch 16/25

60000/60000 [=====] - 3s 52us/step - loss: 0.0130 - acc: 0.9966 - val.

Epoch 17/25

60000/60000 [=====] - 3s 53us/step - loss: 0.0110 - acc: 0.9970 - val.

Epoch 18/25

60000/60000 [=====] - 3s 53us/step - loss: 0.0169 - acc: 0.9951 - val.

Epoch 19/25

60000/60000 [=====] - 3s 52us/step - loss: 0.0090 - acc: 0.9973 - val.

Epoch 20/25

60000/60000 [=====] - 3s 52us/step - loss: 0.0137 - acc: 0.9965 - val.


Epoch 21/25

60000/60000 [=====] - 3s 52us/step - loss: 0.0143 - acc: 0.9960 - val.

Epoch 22/25

60000/60000 [=====] - 3s 52us/step - loss: 0.0103 - acc: 0.9974 - val.

Epoch 23/25
60000/60000 [=====] - 3s 52us/step - loss: 0.0088 - acc: 0.9975 - val.
Epoch 24/25
60000/60000 [=====] - 3s 52us/step - loss: 0.0088 - acc: 0.9976 - val.
Epoch 25/25
60000/60000 [=====] - 3s 53us/step - loss: 0.0095 - acc: 0.9972 - val.
Test score: 0.13170449254391386
Test accuracy: 0.9788



4 Defining a function for model with dropout and no batch normalisation.

In [0]: # <https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization->

```
def nn3(n_layers,d):
    model = Sequential()

    model.add(Dense(n_layers[0], activation='relu', input_shape=(input_dim,), kernel_initializer=K.he_normal(seed=1)))
    model.add(Dropout(d[0]))

    for i in range(len(n_layers)-1):
        model.add(Dense(n_layers[i+1], activation='relu', kernel_initializer=K.he_normal(seed=1)))
```

```

model.add(Dropout(d[i+1]))

model.add(Dense(output_dim, activation='softmax'))

model.summary()

model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)

%matplotlib inline
score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

#dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

5 Architecture with 3 hidden layers.

```

In [81]: n_layers = [512,256,128]
         d = [0.5,0.6,0.5]
         nb_epoch = 25
         nn3(n_layers,d)

```

Layer (type)

Output Shape

Param #

dense_135 (Dense)	(None, 512)	401920
dropout_70 (Dropout)	(None, 512)	0
dense_136 (Dense)	(None, 256)	131328
dropout_71 (Dropout)	(None, 256)	0
dense_137 (Dense)	(None, 128)	32896
dropout_72 (Dropout)	(None, 128)	0
dense_138 (Dense)	(None, 10)	1290

Total params: 567,434

Trainable params: 567,434

Non-trainable params: 0

Train on 60000 samples, validate on 10000 samples

Epoch 1/25

60000/60000 [=====] - 13s 211us/step - loss: 0.6933 - acc: 0.7798 - val.

Epoch 2/25

60000/60000 [=====] - 3s 55us/step - loss: 0.2783 - acc: 0.9225 - val.

Epoch 3/25

60000/60000 [=====] - 3s 53us/step - loss: 0.2091 - acc: 0.9422 - val.

Epoch 4/25

60000/60000 [=====] - 3s 51us/step - loss: 0.1721 - acc: 0.9521 - val.

Epoch 5/25

60000/60000 [=====] - 3s 51us/step - loss: 0.1572 - acc: 0.9570 - val.

Epoch 6/25

60000/60000 [=====] - 3s 51us/step - loss: 0.1416 - acc: 0.9613 - val.

Epoch 7/25

60000/60000 [=====] - 3s 52us/step - loss: 0.1299 - acc: 0.9645 - val.

Epoch 8/25

60000/60000 [=====] - 3s 52us/step - loss: 0.1183 - acc: 0.9675 - val.

Epoch 9/25

60000/60000 [=====] - 3s 52us/step - loss: 0.1090 - acc: 0.9695 - val.

Epoch 10/25

60000/60000 [=====] - 3s 51us/step - loss: 0.1031 - acc: 0.9710 - val.

Epoch 11/25

60000/60000 [=====] - 3s 51us/step - loss: 0.1015 - acc: 0.9721 - val.

Epoch 12/25

60000/60000 [=====] - 3s 52us/step - loss: 0.0935 - acc: 0.9738 - val.


Epoch 13/25

60000/60000 [=====] - 3s 52us/step - loss: 0.0904 - acc: 0.9743 - val.

Epoch 14/25

60000/60000 [=====] - 3s 52us/step - loss: 0.0874 - acc: 0.9759 - val.

Epoch 15/25
60000/60000 [=====] - 3s 51us/step - loss: 0.0850 - acc: 0.9758 - val.
Epoch 16/25
60000/60000 [=====] - 3s 52us/step - loss: 0.0830 - acc: 0.9765 - val.
Epoch 17/25
60000/60000 [=====] - 3s 51us/step - loss: 0.0775 - acc: 0.9776 - val.
Epoch 18/25
60000/60000 [=====] - 3s 51us/step - loss: 0.0771 - acc: 0.9788 - val.
Epoch 19/25
60000/60000 [=====] - 3s 51us/step - loss: 0.0742 - acc: 0.9796 - val.
Epoch 20/25
60000/60000 [=====] - 3s 50us/step - loss: 0.0693 - acc: 0.9801 - val.
Epoch 21/25
60000/60000 [=====] - 3s 50us/step - loss: 0.0712 - acc: 0.9803 - val.
Epoch 22/25
60000/60000 [=====] - 3s 51us/step - loss: 0.0669 - acc: 0.9804 - val.
Epoch 23/25
60000/60000 [=====] - 3s 51us/step - loss: 0.0665 - acc: 0.9805 - val.
Epoch 24/25
60000/60000 [=====] - 3s 51us/step - loss: 0.0622 - acc: 0.9822 - val.
Epoch 25/25
60000/60000 [=====] - 3s 51us/step - loss: 0.0605 - acc: 0.9829 - val.
Test score: 0.07005962141250102
Test accuracy: 0.9826



6 Architecture with 2 hidden layers.

```
In [82]: n_layers = [468,92]
         d = [0.6,0.5]
         nb_epoch = 25
         nn3(n_layers,d)
```

Layer (type)	Output Shape	Param #
dense_139 (Dense)	(None, 468)	367380
dropout_73 (Dropout)	(None, 468)	0
dense_140 (Dense)	(None, 92)	43148
dropout_74 (Dropout)	(None, 92)	0
dense_141 (Dense)	(None, 10)	930

```
=====
Total params: 411,458
Trainable params: 411,458
Non-trainable params: 0
```

```
-----
Train on 60000 samples, validate on 10000 samples
```

```
Epoch 1/25
```

```
60000/60000 [=====] - 12s 204us/step - loss: 0.5460 - acc: 0.8301 - val.
```

```
Epoch 2/25
```

```
60000/60000 [=====] - 3s 48us/step - loss: 0.2620 - acc: 0.9256 - val.
```

```
Epoch 3/25
```

```
60000/60000 [=====] - 3s 48us/step - loss: 0.1996 - acc: 0.9425 - val.
```

```
Epoch 4/25
```

```
60000/60000 [=====] - 3s 47us/step - loss: 0.1770 - acc: 0.9495 - val.
```

```
Epoch 5/25
```

```
60000/60000 [=====] - 3s 47us/step - loss: 0.1545 - acc: 0.9549 - val.
```

```
Epoch 6/25
```

```
60000/60000 [=====] - 3s 48us/step - loss: 0.1453 - acc: 0.9576 - val.
```

```
Epoch 7/25
```

```
60000/60000 [=====] - 3s 48us/step - loss: 0.1310 - acc: 0.9618 - val.
```

```
Epoch 8/25
```

```
60000/60000 [=====] - 3s 48us/step - loss: 0.1191 - acc: 0.9645 - val.
```

```
Epoch 9/25
```

```
60000/60000 [=====] - 3s 47us/step - loss: 0.1163 - acc: 0.9656 - val.
```

```
Epoch 10/25
```

```
60000/60000 [=====] - 3s 47us/step - loss: 0.1085 - acc: 0.9682 - val.
```

```
Epoch 11/25
```


```
60000/60000 [=====] - 3s 47us/step - loss: 0.1048 - acc: 0.9691 - val.
```

```
Epoch 12/25
```

```

60000/60000 [=====] - 3s 47us/step - loss: 0.0997 - acc: 0.9707 - val.
Epoch 13/25
60000/60000 [=====] - 3s 48us/step - loss: 0.0954 - acc: 0.9715 - val.
Epoch 14/25
60000/60000 [=====] - 3s 47us/step - loss: 0.0937 - acc: 0.9730 - val.
Epoch 15/25
60000/60000 [=====] - 3s 47us/step - loss: 0.0901 - acc: 0.9732 - val.
Epoch 16/25
60000/60000 [=====] - 3s 47us/step - loss: 0.0835 - acc: 0.9750 - val.
Epoch 17/25
60000/60000 [=====] - 3s 48us/step - loss: 0.0852 - acc: 0.9744 - val.
Epoch 18/25
60000/60000 [=====] - 3s 47us/step - loss: 0.0828 - acc: 0.9748 - val.
Epoch 19/25
60000/60000 [=====] - 3s 47us/step - loss: 0.0802 - acc: 0.9758 - val.
Epoch 20/25
60000/60000 [=====] - 3s 47us/step - loss: 0.0769 - acc: 0.9768 - val.
Epoch 21/25
60000/60000 [=====] - 3s 48us/step - loss: 0.0743 - acc: 0.9778 - val.
Epoch 22/25
60000/60000 [=====] - 3s 48us/step - loss: 0.0767 - acc: 0.9764 - val.
Epoch 23/25
60000/60000 [=====] - 3s 48us/step - loss: 0.0728 - acc: 0.9777 - val.
Epoch 24/25
60000/60000 [=====] - 3s 47us/step - loss: 0.0712 - acc: 0.9783 - val.
Epoch 25/25
60000/60000 [=====] - 3s 47us/step - loss: 0.0681 - acc: 0.9795 - val.
Test score: 0.07291174245599469
Test accuracy: 0.9823

```



7 Architecture with 5 hidden layers

```
In [83]: n_layers = [650,520,410,290,125]
         d = [0.5,0.6,0.5,0.5,0.5]
         nb_epoch = 25
         nn3(n_layers,d)
```

Layer (type)	Output Shape	Param #
dense_142 (Dense)	(None, 650)	510250
dropout_75 (Dropout)	(None, 650)	0
dense_143 (Dense)	(None, 520)	338520
dropout_76 (Dropout)	(None, 520)	0
dense_144 (Dense)	(None, 410)	213610
dropout_77 (Dropout)	(None, 410)	0
dense_145 (Dense)	(None, 290)	119190

dropout_78 (Dropout)	(None, 290)	0

dense_146 (Dense)	(None, 125)	36375

dropout_79 (Dropout)	(None, 125)	0


dense_147 (Dense)	(None, 10)	1260
=====		
Total params: 1,219,205		
Trainable params: 1,219,205		
Non-trainable params: 0		

Train on 60000 samples, validate on 10000 samples		
Epoch 1/25		
60000/60000 [=====] - 13s 217us/step - loss: 1.1927 - acc: 0.5892 - val.		
Epoch 2/25		
60000/60000 [=====] - 4s 61us/step - loss: 0.3680 - acc: 0.9039 - val.		
Epoch 3/25		
60000/60000 [=====] - 4s 60us/step - loss: 0.2690 - acc: 0.9318 - val.		
Epoch 4/25		
60000/60000 [=====] - 4s 60us/step - loss: 0.2314 - acc: 0.9419 - val.		
Epoch 5/25		
60000/60000 [=====] - 4s 59us/step - loss: 0.2028 - acc: 0.9509 - val.		
Epoch 6/25		
60000/60000 [=====] - 4s 59us/step - loss: 0.1827 - acc: 0.9546 - val.		
Epoch 7/25		
60000/60000 [=====] - 4s 61us/step - loss: 0.1654 - acc: 0.9595 - val.		
Epoch 8/25		
60000/60000 [=====] - 4s 61us/step - loss: 0.1588 - acc: 0.9606 - val.		
Epoch 9/25		
60000/60000 [=====] - 4s 62us/step - loss: 0.1512 - acc: 0.9623 - val.		
Epoch 10/25		
60000/60000 [=====] - 4s 61us/step - loss: 0.1447 - acc: 0.9647 - val.		
Epoch 11/25		
60000/60000 [=====] - 4s 60us/step - loss: 0.1356 - acc: 0.9666 - val.		
Epoch 12/25		
60000/60000 [=====] - 4s 59us/step - loss: 0.1276 - acc: 0.9689 - val.		
Epoch 13/25		
60000/60000 [=====] - 4s 60us/step - loss: 0.1289 - acc: 0.9686 - val.		
Epoch 14/25		
60000/60000 [=====] - 4s 60us/step - loss: 0.1214 - acc: 0.9703 - val.		
Epoch 15/25		
60000/60000 [=====] - 4s 60us/step - loss: 0.1139 - acc: 0.9728 - val.		
Epoch 16/25		
60000/60000 [=====] - 4s 59us/step - loss: 0.1100 - acc: 0.9727 - val.		
Epoch 17/25		
60000/60000 [=====] - 4s 60us/step - loss: 0.1101 - acc: 0.9729 - val.		
Epoch 18/25		


```

60000/60000 [=====] - 4s 59us/step - loss: 0.1053 - acc: 0.9743 - val.
Epoch 19/25
60000/60000 [=====] - 4s 60us/step - loss: 0.1106 - acc: 0.9731 - val.
Epoch 20/25
60000/60000 [=====] - 4s 60us/step - loss: 0.1019 - acc: 0.9760 - val.
Epoch 21/25
60000/60000 [=====] - 4s 60us/step - loss: 0.0978 - acc: 0.9759 - val.
Epoch 22/25
60000/60000 [=====] - 4s 60us/step - loss: 0.0950 - acc: 0.9769 - val.
Epoch 23/25
60000/60000 [=====] - 4s 60us/step - loss: 0.0943 - acc: 0.9773 - val.
Epoch 24/25
60000/60000 [=====] - 4s 60us/step - loss: 0.0916 - acc: 0.9777 - val.
Epoch 25/25
60000/60000 [=====] - 4s 59us/step - loss: 0.0877 - acc: 0.9781 - val.
Test score: 0.08099026859523369
Test accuracy: 0.9819

```



8 Experimenting with conditional batch normalisation and dropout with 3 hidden layers.

```

In [0]: # https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-
def expt():

```

```

model = Sequential()
model.add(Dense(584, activation='relu', input_shape=(input_dim,), kernel_initializer=
model.add(BatchNormalization())
model.add(Dropout(0.5))
model.add(Dense(320, activation='relu',kernel_initializer=K.he_normal(seed=None)) )
model.add(BatchNormalization())
model.add(Dense(125, activation='relu',kernel_initializer=K.he_normal(seed=None)) )
model.add(Dropout(0.5))

model.add(Dense(output_dim, activation='softmax'))

model.summary()

model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)

%matplotlib inline
score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

#dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

In [87]: expt()

Layer (type)	Output Shape	Param #
dense_148 (Dense)	(None, 584)	458440
batch_normalization_90 (Batch Normalization)	(None, 584)	2336
dropout_80 (Dropout)	(None, 584)	0
dense_149 (Dense)	(None, 320)	187200
batch_normalization_91 (Batch Normalization)	(None, 320)	1280
dense_150 (Dense)	(None, 125)	40125
dropout_81 (Dropout)	(None, 125)	0
dense_151 (Dense)	(None, 10)	1260

Total params: 690,641

Trainable params: 688,833

Non-trainable params: 1,808

Train on 60000 samples, validate on 10000 samples

Epoch 1/25

60000/60000 [=====] - 14s 240us/step - loss: 0.4274 - acc: 0.8725 - val_loss: 0.4274 - val_acc: 0.8725

Epoch 2/25

60000/60000 [=====] - 4s 73us/step - loss: 0.1952 - acc: 0.9413 - val_loss: 0.1952 - val_acc: 0.9413

Epoch 3/25

60000/60000 [=====] - 4s 73us/step - loss: 0.1537 - acc: 0.9543 - val_loss: 0.1537 - val_acc: 0.9543

Epoch 4/25

60000/60000 [=====] - 4s 75us/step - loss: 0.1326 - acc: 0.9605 - val_loss: 0.1326 - val_acc: 0.9605

Epoch 5/25

60000/60000 [=====] - 4s 74us/step - loss: 0.1166 - acc: 0.9649 - val_loss: 0.1166 - val_acc: 0.9649

Epoch 6/25

60000/60000 [=====] - 5s 75us/step - loss: 0.1027 - acc: 0.9687 - val_loss: 0.1027 - val_acc: 0.9687

Epoch 7/25

60000/60000 [=====] - 5s 76us/step - loss: 0.0979 - acc: 0.9701 - val_loss: 0.0979 - val_acc: 0.9701

Epoch 8/25

60000/60000 [=====] - 5s 76us/step - loss: 0.0907 - acc: 0.9718 - val_loss: 0.0907 - val_acc: 0.9718

Epoch 9/25

60000/60000 [=====] - 4s 72us/step - loss: 0.0847 - acc: 0.9740 - val_loss: 0.0847 - val_acc: 0.9740

Epoch 10/25

60000/60000 [=====] - 4s 73us/step - loss: 0.0812 - acc: 0.9745 - val_loss: 0.0812 - val_acc: 0.9745

Epoch 11/25

60000/60000 [=====] - 4s 75us/step - loss: 0.0734 - acc: 0.9773 - val_loss: 0.0734 - val_acc: 0.9773

Epoch 12/25


60000/60000 [=====] - 4s 71us/step - loss: 0.0717 - acc: 0.9780 - val_loss: 0.0717 - val_acc: 0.9780

Epoch 13/25

```

60000/60000 [=====] - 4s 72us/step - loss: 0.0663 - acc: 0.9793 - val.
Epoch 14/25
60000/60000 [=====] - 4s 72us/step - loss: 0.0651 - acc: 0.9792 - val.
Epoch 15/25
60000/60000 [=====] - 4s 72us/step - loss: 0.0624 - acc: 0.9805 - val.
Epoch 16/25
60000/60000 [=====] - 4s 73us/step - loss: 0.0586 - acc: 0.9817 - val.
Epoch 17/25
60000/60000 [=====] - 4s 72us/step - loss: 0.0573 - acc: 0.9821 - val.
Epoch 18/25
60000/60000 [=====] - 4s 72us/step - loss: 0.0550 - acc: 0.9823 - val.
Epoch 19/25
60000/60000 [=====] - 4s 71us/step - loss: 0.0526 - acc: 0.9840 - val.
Epoch 20/25
60000/60000 [=====] - 4s 71us/step - loss: 0.0478 - acc: 0.9846 - val.
Epoch 21/25
60000/60000 [=====] - 4s 70us/step - loss: 0.0479 - acc: 0.9845 - val.
Epoch 22/25
60000/60000 [=====] - 4s 71us/step - loss: 0.0469 - acc: 0.9852 - val.
Epoch 23/25
60000/60000 [=====] - 4s 73us/step - loss: 0.0453 - acc: 0.9855 - val.
Epoch 24/25
60000/60000 [=====] - 4s 72us/step - loss: 0.0473 - acc: 0.9850 - val.
Epoch 25/25
60000/60000 [=====] - 4s 71us/step - loss: 0.0415 - acc: 0.9864 - val.
Test score: 0.05252337242660105
Test accuracy: 0.9855

```



9 Conclusions

```
In [90]: from prettytable import PrettyTable
print("=====")
print(" FOR DROPOUT+BATCH NORMALISATION")
x = PrettyTable()
x.field_names = ["No. of hidden layers", "Train Loss", " Val Loss", "Val AUC"]
x.add_row([2, 0.0645, 0.0584, 0.9828])
x.add_row([3, 0.0612, 0.0549, 0.9850])
x.add_row([4, 0.0612, 0.0638, 0.9847])
x.add_row([5, 0.0669, 0.0607, 0.9846])
print(x)
print("=====")
print("=====")
print(" FOR ONLY BATCH NORMALISATION and NO DROPOUT")
x = PrettyTable()
x.field_names = ["No. of hidden layers", "Train Loss", " Val Loss", "Val AUC"]
x.add_row([2, 0.0050, 0.0840, 0.9808])
x.add_row([3, 0.0058, 0.0723, 0.9833])
x.add_row([5, 0.0093, 0.0698, 0.9840])
print(x)

print("=====")
print("=====")
```

```

print(" FOR NO DROPOUT AND NO BATCH NORMALISATION")
x = PrettyTable()
x.field_names = ["No. of hidden layers", "Train Loss", " Val Loss", "Val AUC"]
x.add_row([2, 0.0065, 0.0998, 0.9807])
x.add_row([3, 0.0072, 0.0933, 0.9814])
x.add_row([5, 0.0095, 0.1317, 0.9788])
print(x)

print("=====")
print("=====")
print(" FOR DROPOUT ONLY AND NO BATCH NORMALISATION")
x = PrettyTable()
x.field_names = ["No. of hidden layers", "Train Loss", " Val Loss", "Val AUC"]
x.add_row([2, 0.0681, 0.0729, 0.9823])
x.add_row([3, 0.0605, 0.0701, 0.9826])
x.add_row([5, 0.0877, 0.0810, 0.9819])
print(x)

print("=====")
print("=====")
print("EXPERIMENTAL _ DROPOUT-NORMALISATION-DROPOUT")
x = PrettyTable()
x.field_names = ["No. of hidden layers", "Train Loss", " Val Loss", "Val AUC"]
x.add_row([3, 0.0550, 0.0556, 0.9845])
print(x)

```

```

=====
FOR DROPOUT+BATCH NORMALISATION
+-----+-----+-----+-----+
| No. of hidden layers | Train Loss | Val Loss | Val AUC |
+-----+-----+-----+-----+
|          2          |    0.0645   |    0.0584   |    0.9828 |
|          3          |    0.0612   |    0.0549   |    0.985   |
|          4          |    0.0612   |    0.0638   |    0.9847 |
|          5          |    0.0669   |    0.0607   |    0.9846 |
+-----+-----+-----+-----+
=====
=====
FOR ONLY BATCH NORMALISATION and NO DROPOUT
+-----+-----+-----+-----+
| No. of hidden layers | Train Loss | Val Loss | Val AUC |
+-----+-----+-----+-----+
|          2          |    0.005    |    0.084    |    0.9808 |
|          3          |    0.0058   |    0.0723   |    0.9833 |
|          5          |    0.0093   |    0.0698   |    0.984   |
+-----+-----+-----+-----+
=====
=====

```

FOR NO DROPOUT AND NO BATCH NORMALISATION

No. of hidden layers	Train Loss	Val Loss	Val AUC
2	0.0065	0.0998	0.9807
3	0.0072	0.0933	0.9814
5	0.0095	0.1317	0.9788

FOR DROPOUT ONLY AND NO BATCH NORMALISATION

No. of hidden layers	Train Loss	Val Loss	Val AUC
2	0.0681	0.0729	0.9823
3	0.0605	0.0701	0.9826
5	0.0877	0.081	0.9819

EXPERIMENTAL _ DROPOUT-NORMALISATION-DROPOUT

No. of hidden layers	Train Loss	Val Loss	Val AUC
3	0.055	0.0556	0.9845

9.0.1 Conclusions

- 1) Observing all the pretty tables, we conclude that (Dropout+Batch Normalisation) > only Dropout > only batch normalisation > no dropout or normalisation
- 2) As can be observed from the graphs , overfitting occurs, wherever there is no dropout or batch normalisation.
- 3) The best model was found from the experimental model.(Considering upto epoch 18.) It was done as : For this , drop out was applied on first hidden layer.Then, batch normalisation on second hidden layer. Then again, dropout on third hidden layer. This way,instead of applying both of them on all layers, time and cost of computing was saved. The result was also good. The diff between train loss and val loss was 0.0006 at epoch 18(Least of all models).Other models took till epoch near to 25 to converge.
- 4) We can infer that dropout and batch normalisation helps a lot in deep learning models especially in avoiding overfitting.